
Chapter 3: Processes and Inter Process Communications

Processes

❑ Objectives

- To introduce the notion of a process – a program in execution, which forms the basis of all computation
- To describe the various features of processes, including scheduling, creation and termination, and communication


❑ Topics

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems

Process Concept

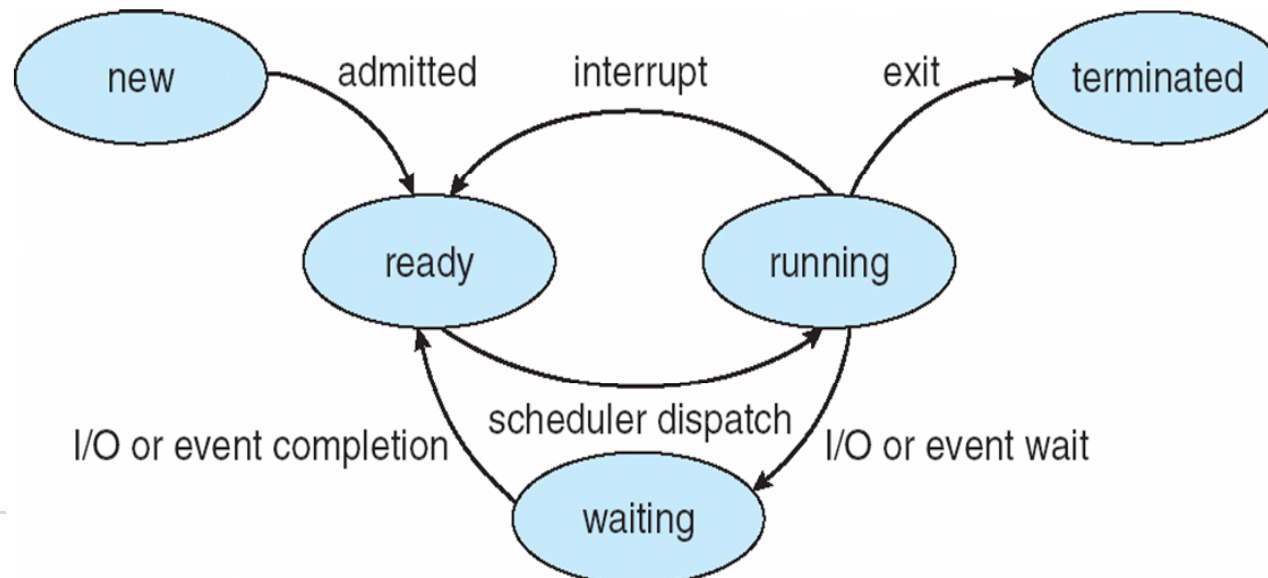
- ❑ An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- ❑ Textbook uses the terms job and process almost interchangeably
- ❑ Process – a program in execution; process execution must progress in a sequential fashion
- ❑ A process includes:
 - Program counter
 - Stack (and heap)
 - Data section

Process in Memory



Process States

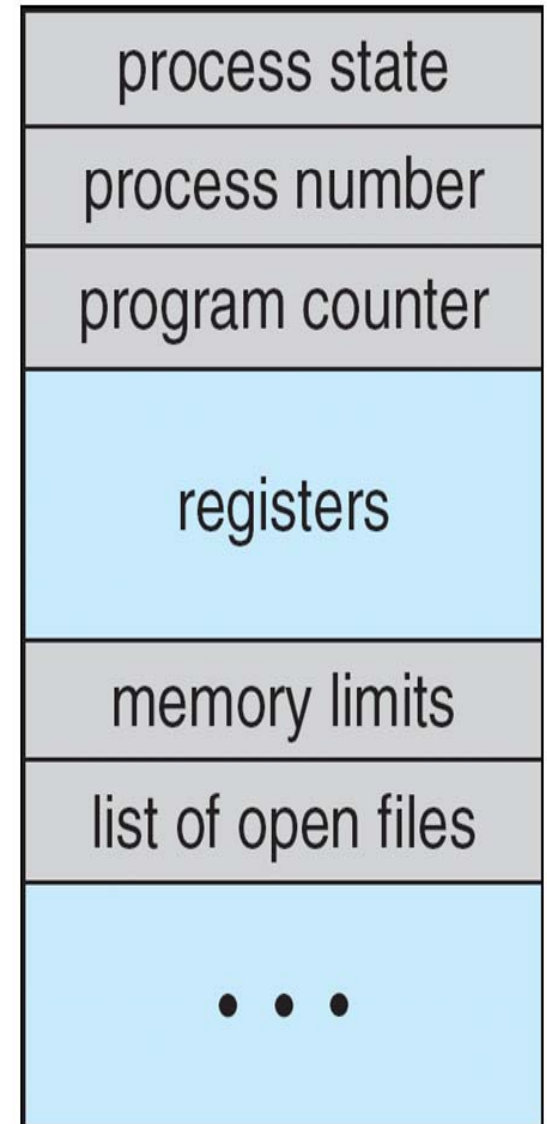
- ❑ As a process executes, it changes state (Linux)
 - new: The process is being created
 - running: Instructions are being executed
 - waiting: The process is waiting for some event to occur
 - ready: The process is waiting to be assigned to a processor
 - terminated: The process has finished execution



Process Control Block (PCB)

❑ Information associated with each process

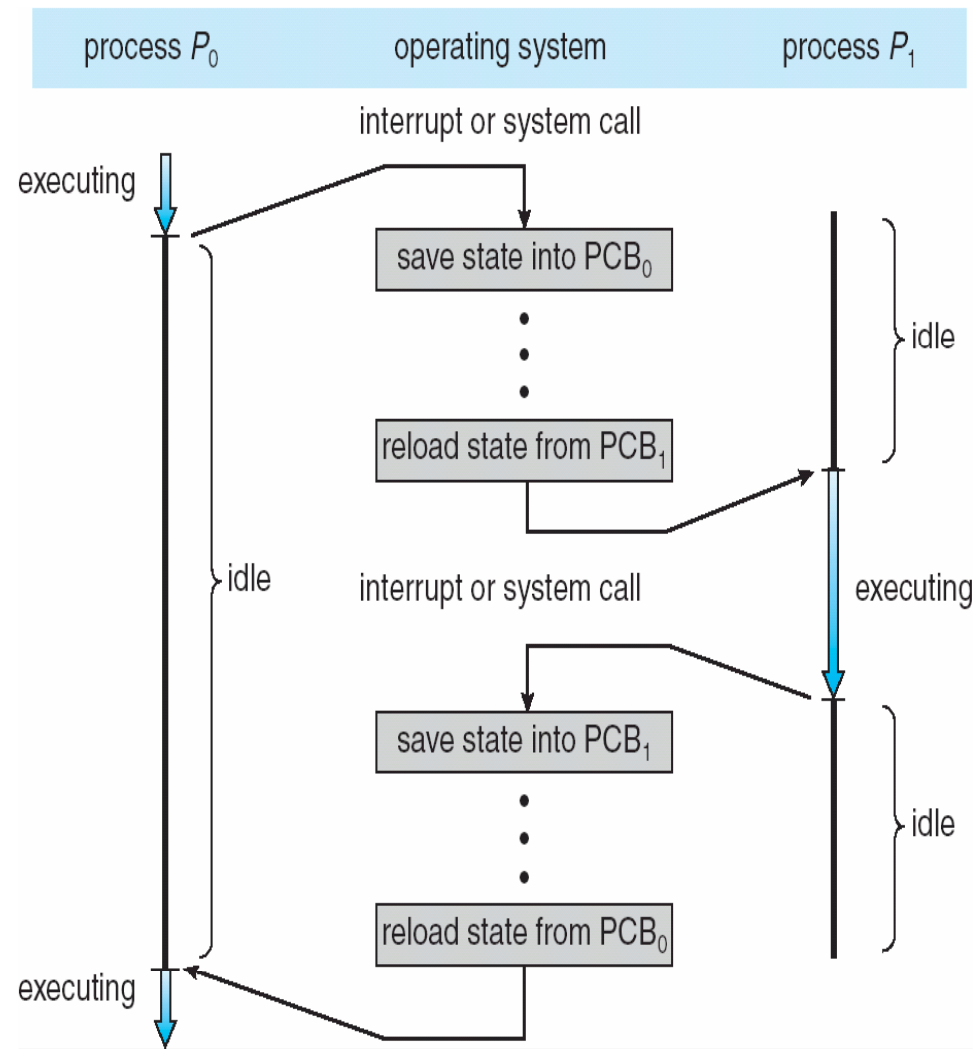
- Process state
 - Process number
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information



Context Switch

- ❑ When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a context switch
- ❑ Context of a process represented in the PCB
- ❑ Context-switch time is overhead
 - The system does no useful work while switching
- ❑ Time is dependent on hardware support

CPU Switch From Process to Process

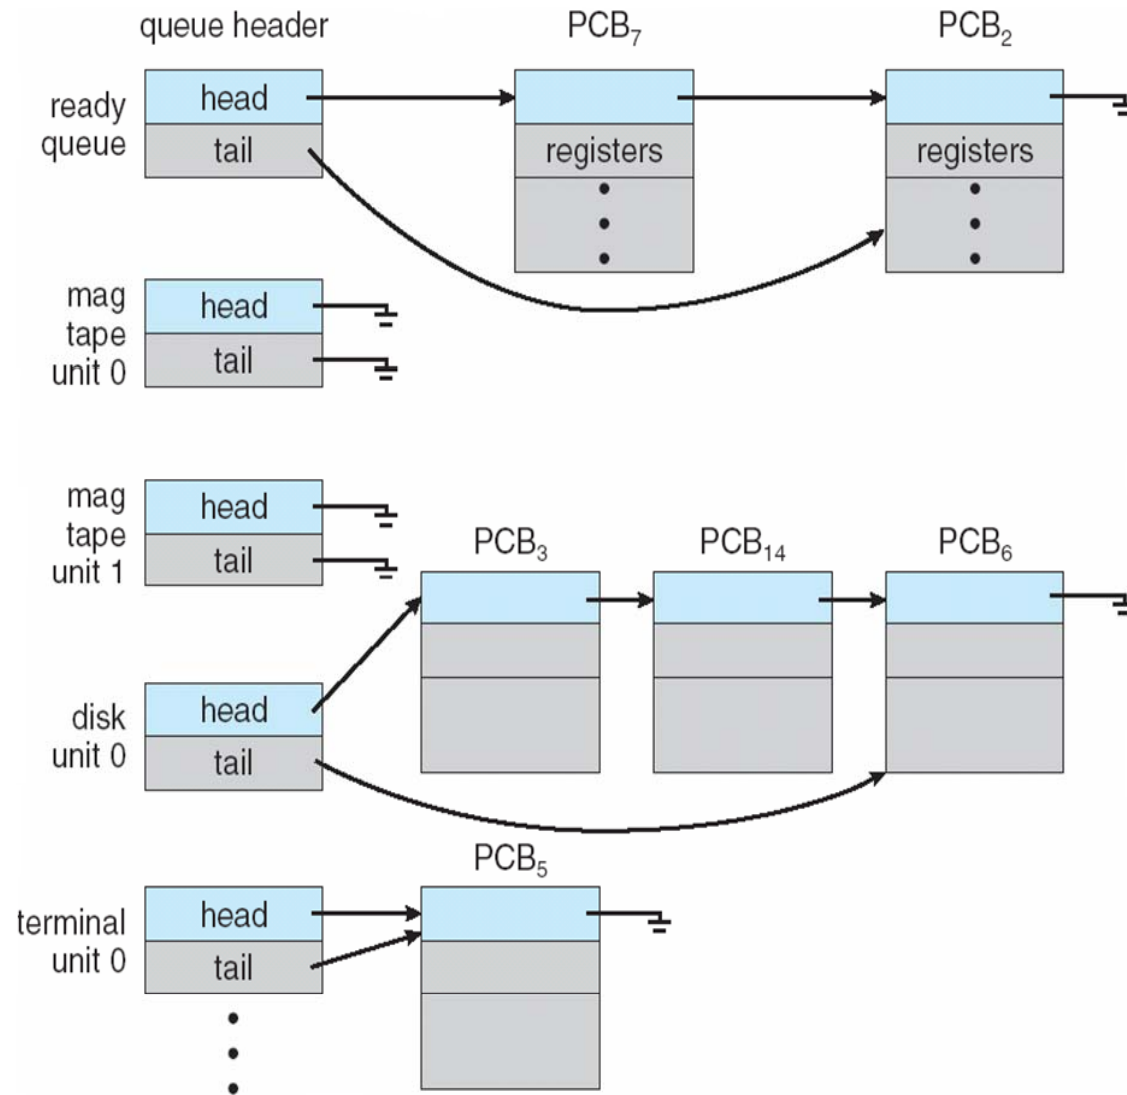


Process Scheduling Queues

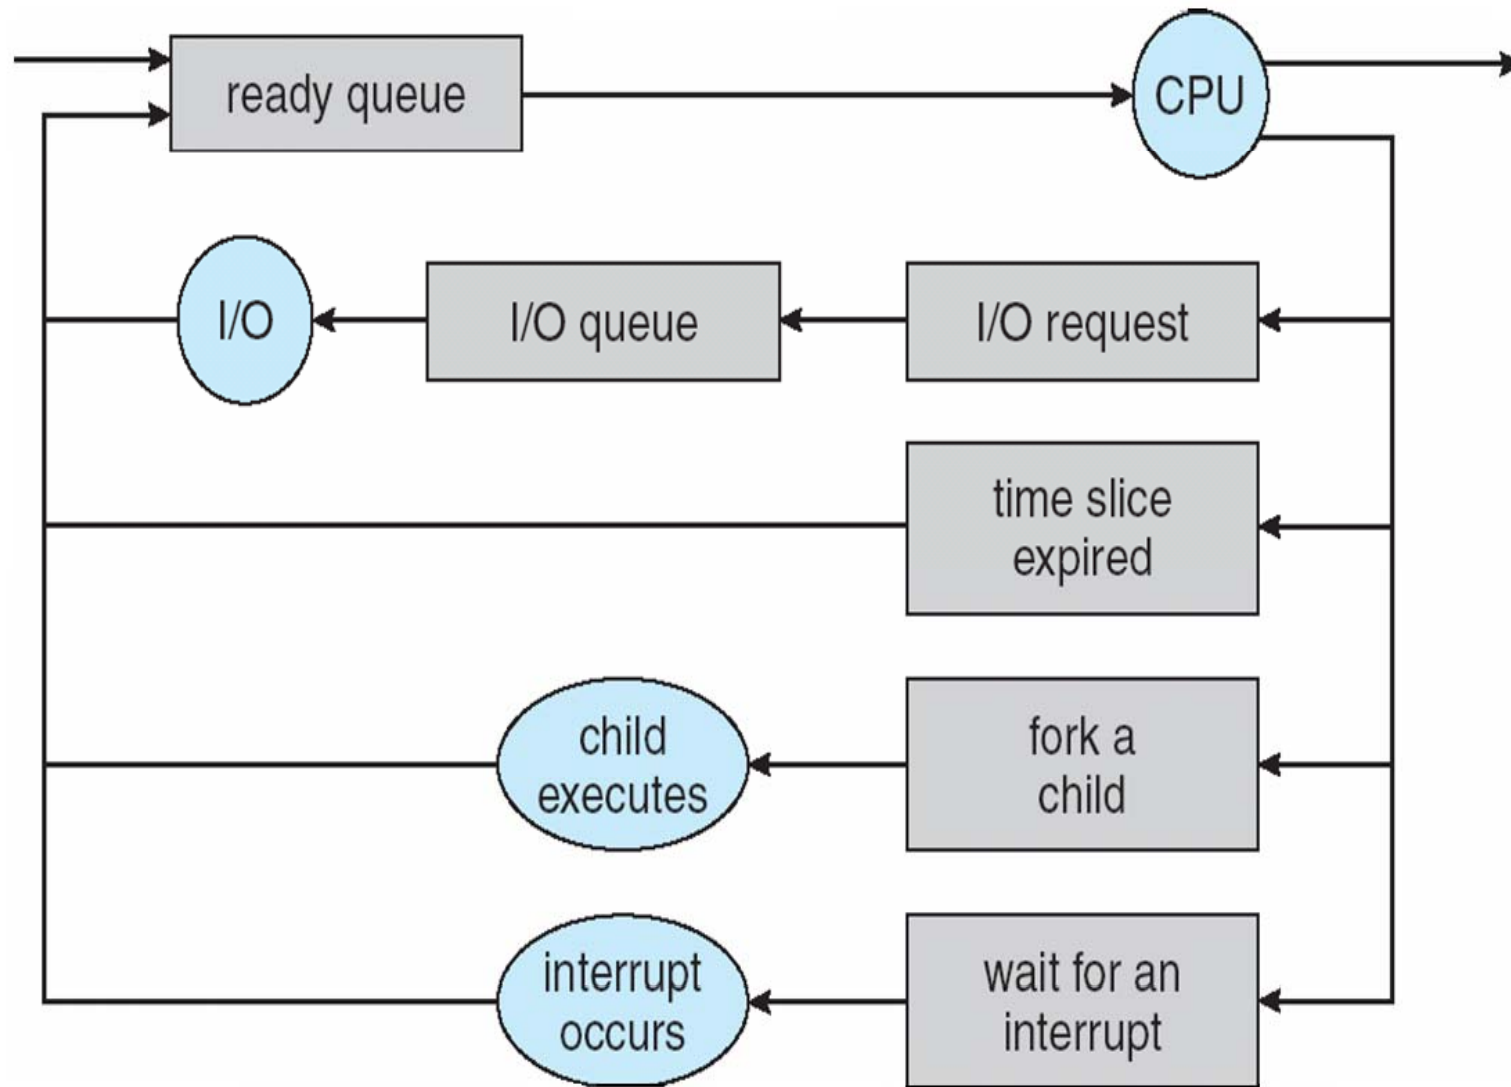
- ❑ Job queue
 - Set of all processes in the system
- ❑ Ready queue
 - Set of all processes residing in main memory, ready and waiting to execute
- ❑ Device queues
 - Set of processes waiting for an I/O device

- ❑ Processes migrate among the various queues

Ready Queue and Various I/O Device Queues

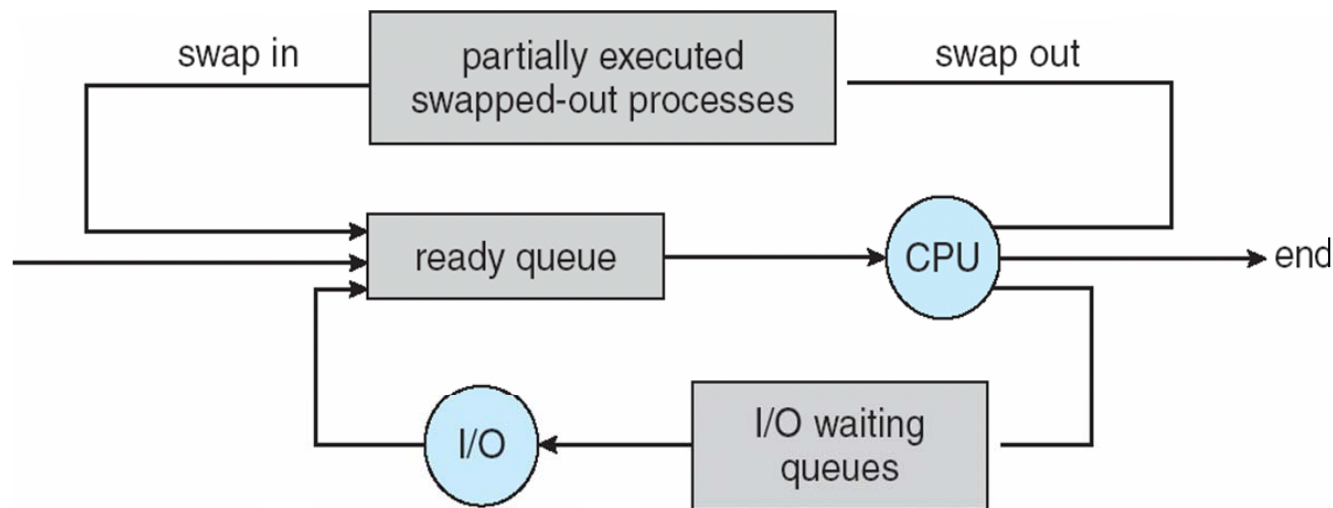


Representation of Process Scheduling



Schedulers (1)

- ❑ Long-term scheduler (or job scheduler)
 - Selects which processes should be brought into the ready queue
- ❑ Short-term scheduler (or CPU scheduler)
 - Selects which process should be executed next and allocates CPU
- ❑ Addition of medium-term scheduling
 - Swapping process in/out



Schedulers (2)

- ❑ Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)
- ❑ Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
 - May be absent, jobs simply added to ready queue
- ❑ Processes can be described as either
 - I/O-bound process – spends more time doing I/O than computations, many short CPU bursts
 - CPU-bound process – spends more time doing computations; few very long CPU bursts

Process Creation (1)

- ❑ Parent process creates children processes, creates other processes, which results in forming a tree of processes
- ❑ Generally, all processes are identified and managed via a process identifier (pid)
- ❑ Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources

Process Creation (2)

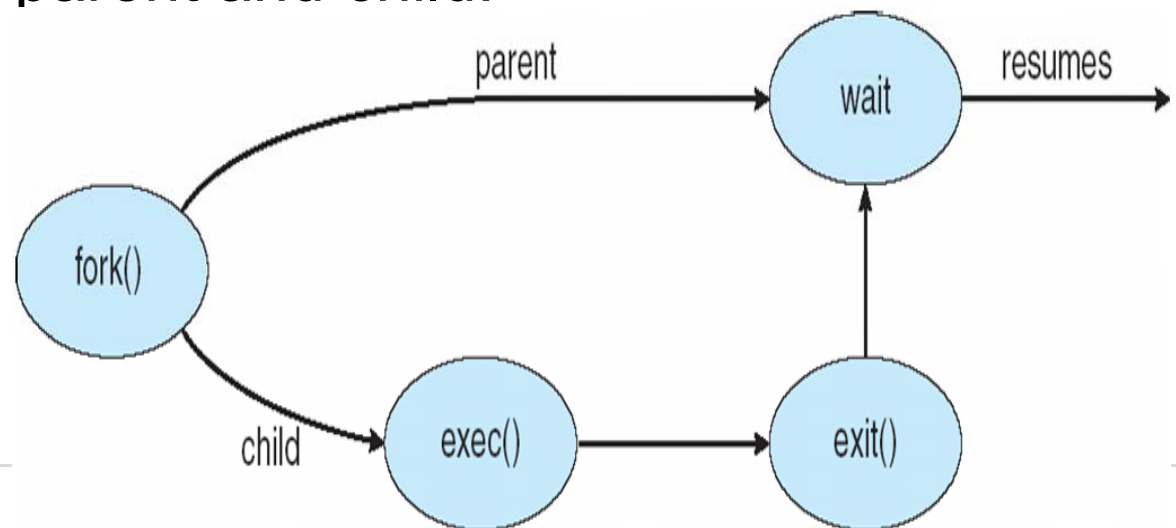
- ❑ Execution strategy options
 - Parent and children execute concurrently
 - Parent waits until children terminate

- ❑ Address space options
 - Child is a duplicate of parent
 - Child has a program loaded into it

Explicit Process Creation

- ❑ UNIX examples
 - fork system call creates new process
- ❑ `procid = fork()` replicates calling process
 - `exec` system call used after a fork to replace the process' memory space with a new program
- ❑ Parent and child identical except for the value of `procid`
 - Use `procid` to diverge parent and child:

```
if (procid == 0)
    do_child_processing
else
    do_parent_processing
```



C Program Forking Separate Process

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

int main()
{
    pid_t pid;

    pid = fork(); /* fork a child process */

    if (pid < 0) { /* error occurred */
        printf("Fork Failed\n");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        printf("I am the child %d\n", pid);
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        printf("I am the parent %d\n", pid); /* parent will wait for the child to complete */
        wait(NULL);


        printf("Child Complete\n");
        exit(0);
    }
}
```

Process Termination

- ❑ Process executes last statement and asks the operating system to delete it (exit)
 - Output data from child to parent (via wait)
 - Process' resources are deallocated by operating system

- ❑ Parent may terminate execution of children processes
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - Some operating system do not allow child to continue if its parent terminates
 - All children terminated - cascading termination
 - Linux: reparenting to pid 1 (init)

Process Ordering and Precedence



- Process flow graph depicts process order, parallelism, and precedence
 - Serial execution is expressed as: $S(p_1, p_2, \dots)$
 - Parallel execution is expressed as: $P(p_1, p_2, \dots)$
- Figure (c) represents the following: $S(p_1, P(p_2, S(p_3, P(p_4, p_5)), p_6), P(p_7, p_8))$

Data Parallelism

- ❑ Same code is applied to different data
 - *E.g.*, SIMD vector instructions

SIMD: Single Instruction Multiple Data

- ❑ The forall statement
 - Syntax: forall (parameters) statements
 - Semantics (meaning):
 - Parameters specify set of data items
 - Statements are executed for each item parallel

Data Parallelism

- Example matrix multiplication: $A = B \times C$

```
forall ( i:1..n, j:1..m )
```

```
    A[i][j] = 0;
```

```
    for ( k=1; k<=r; ++k )
```

```
        A[i][j] = A[i][j] +  
                  B[i][k]*C[k][j];
```


- Each inner product is computed sequentially
- All column-row products are computed in parallel

Interprocess Communication (IPC)

- ❑ Processes in a system: independent or cooperating
 - Independent processes cannot affect or be affected by the execution of another process
 - Cooperating processes can affect or be affected by the execution of another process
- ❑ Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- ❑ Cooperating processes need Interprocess Communication (IPC) mechanisms


Communications Models

- ❑ Two models of IPC
 - Message passing



(a)

Shared memory



(b)

Producer-Consumer Problem

- ❑ Paradigm for cooperating processes, producer process produces information that is consumed by a consumer process
 - Unbounded-buffer places no practical limit on the size of the buffer
 - Bounded-buffer assumes that there is a fixed buffer size
- ❑ Buffer data structure as a circular array of items
 - Using shared memory for IPC

```
#define BUFFER_SIZE 10
class item {
    . . .
};

item
buffer[BUFFER_SIZE];
// points to empty slot
int in = 0;
// next available item
int out = 0;
```


Bounded-Buffer – Producer

```
while (true) {  
    // Produce an item  
    while (((in + 1) % BUFFER SIZE count) ==  
out);  
    // do nothing -- no free buffers  
    buffer[in] = item;  
    in = (in + 1) % BUFFER SIZE;  
}
```

?: remainder after division (modulo division)

Bounded Buffer – Consumer

```
while (true) {  
    while (in == out);  
    // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    // do something with the item  
}
```

Message Passing (1)

- ❑ Mechanism for processes to communicate and to synchronize their actions
- ❑ Message system – processes communicate with each other without resorting to shared variables
- ❑ IPC facility provides two operations (setup – msgget()):
 - send(message) – msgsnd ()
 - receive(message) – msgrcv()

Message Passing (2)

- ❑ If P and Q wish to communicate, they need to
 - Establish a communication link between them
 - Exchange messages via send/receive
- ❑ Implementation of communication link
 - Physical (e.g., shared memory, hardware bus, network)
 - Logical (e.g., logical properties)
- ❑ Logical properties of a link
 - Direct or indirect communication
 - Synchronous or asynchronous communication
 - Automatic or explicit buffering

Direct Communication

- ❑ Processes must name each other explicitly
 - send (P, message) – send a message to process P
 - receive(Q, message) – receive a message from process Q

- ❑ Properties of the communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - receive(id, message) is possible too – receive message from any process (symmetric / asymmetric addressing)

Indirect Communication (1)

- ❑ Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- ❑ Primitives are defined as
 - `send(A, message)` – send a message to mailbox A
 - `receive(A, message)` – receive a message from mailbox A
- ❑ Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links

Indirect Communication (2)

❑ Mailbox sharing

- P1, P2, and P3 share mailbox A
- P1, sends; P2 and P3 receive
- Who gets the message?

❑ Solution

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver
 - Sender is notified who the receiver was

Synchronization

- ❑ Message passing may be either blocking or non-blocking (IPC_NOWAIT)
- ❑ Blocking is considered synchronous
 - Blocking send has the sender block until the message is received
 - Blocking receive has the receiver block until a message is available
- ❑ Non-blocking is considered asynchronous
 - Non-blocking send has the sender send the message and continues
 - Non-blocking receive has the receiver receive a valid message or null

Buffering

- ❑ Queue of messages attached to the link; implemented in one of three ways
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits

Examples of IPC Systems – POSIX

❑ POSIX Shared Memory

- Process first creates shared memory segment
 - `segment id = shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);`
- Process wanting access to that shared memory must attach to it
 - `shared memory = (char *) shmat(id, NULL, 0);`
- Now the process could write to the shared memory
 - `sprintf(shared memory, "Writing to shared memory");`
- When done a process can detach the shared memory from its address space
 - `shmdt(shared memory);`

Examples of IPC Systems – Mach

- ❑ Mach communication is message-based
 - Even system calls are messages
 - Each task gets two mailboxes at creation
 - Kernel and Notify
 - Only three system calls needed for message transfer
 - `msg_send()`, `msg_receive()`, `msg_rpc()`
 - Mailboxes needed for communication, created via
 - `port_allocate()`

Shared Memory + Fork

```
#include <unistd.h>
#include <stdio.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <sys/wait.h>
```

```
int main()
{
```

Shared memory preparation

```
    int pid;
    int segment_id; /* the identifier for the
shared memory segment */
    char *shared_memory; /* a pointer to the
shared memory segment */
    const int segment_size = 4096; /* the size
(in bytes) of the shared memory segment */
```

```
    /** allocate a shared memory segment */
    segment_id = shmget(IPC_PRIVATE,
segment_size, S_IRUSR | S_IWUSR);

    /** attach the shared memory segment */
    shared_memory = (char *) shmat(segment_id,
NULL, 0);
    //printf("shared memory segment %d attached
at address %p\n", segment_id, shared_memory);
```

Shared memory allocation & attachment

```
pid = fork(); /* fork another process */
```

```
    if (pid == 0) { /* child process */
        /** write a message to the shared memory
segment */
        printf("%s\n", shared_memory);
        sprintf(shared_memory, "Thomas was here");
    } else { /* parent process */
        wait(NULL);
```

```
        /** now print out the string from shared
memory */
        printf("*Parent Got: %s*\n",
shared_memory);
```

```
        /** now detach the shared memory segment
*/
        if (shmdt(shared_memory) == -1) {
            fprintf(stderr, "Unable to detach\n");
        }
        /** now remove the shared memory segment
*/
        shmctl(segment_id, IPC_RMID, NULL);

        return 0;
    }
}
```

Forking


Writing

Communications in Client-Server Systems

- ❑ Sockets
- ❑ Remote Procedure Calls
 - Remote Method Invocation (Java)
 - gRPC
 - open source remote procedure call (RPC)
 - Uses HTTP/2, Protocol Buffers
 - Example NFS
- ❑ Pipes

Sockets

- ❑ A socket is defined as an endpoint for communication
- ❑ Concatenation of IP address and port
- ❑ The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
- ❑ Communication consists between a pair of sockets

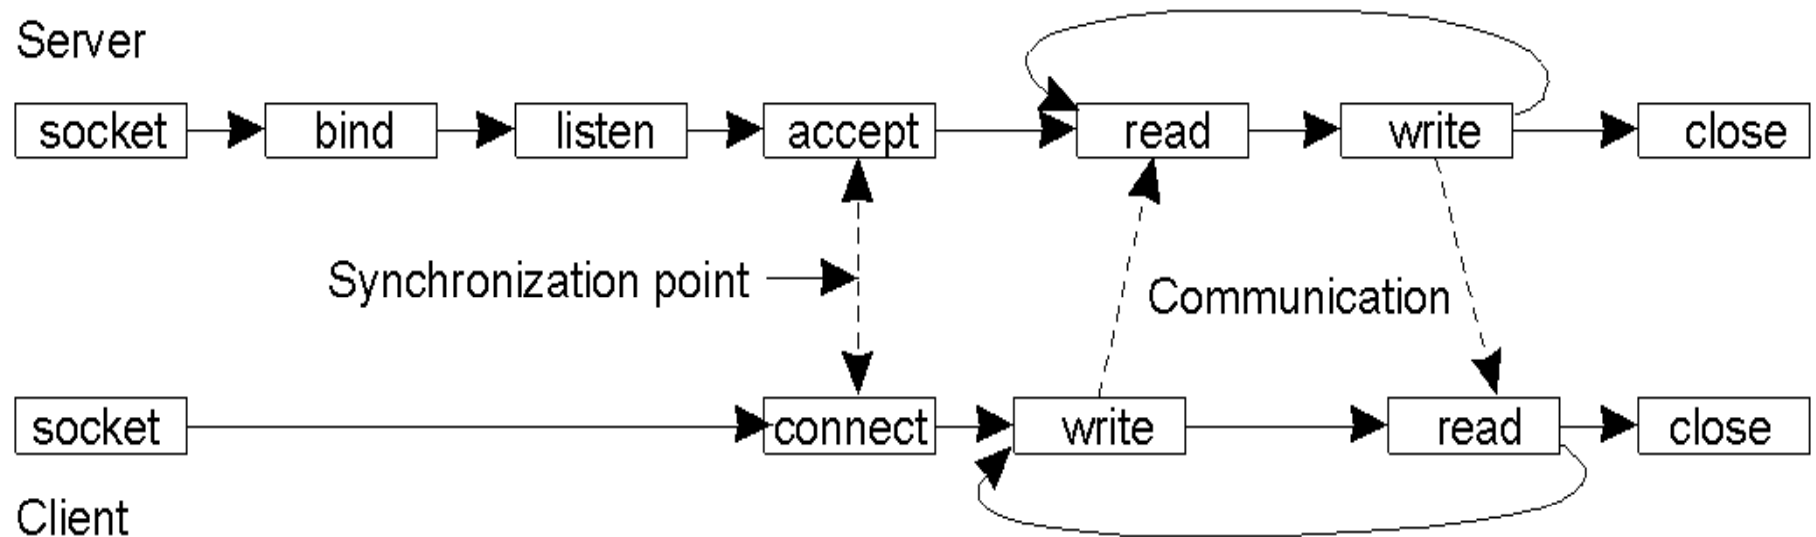


System Calls for TCP/IP Sockets

<i>System call</i>	<i>Who calls it</i>	<i>Meaning</i>
Socket	Server / Client	Create a new communication endpoint
Bind	Server	Attach a local address and port to a socket
Listen	Server	Define how many clients can be queued
Accept	Server	Block until a connection request arrives
Connect	Client	Actively attempt to establish a connection
Write	Server / Client	Send some data over the connection
Read	Server / Client	Receive some data over the connection
Close	Server / Client	Release the connection

TCP/IP Communication

❑ Connection-oriented



Java Example

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(10117);

            // now listen for connections
            while (true) {
                Socket client = sock.accept();
                // we have a connection

                PrintWriter pout = new
PrintWriter(client.getOutputStream(), true);
                // write the Date to the socket
                pout.println(new
java.util.Date().toString());

                // close the socket and resume listening
for more connections
                client.close();
            }
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            // this could be changed to an IP
name or address other than the localhost
            Socket sock = new
Socket("127.0.0.1",10117);
            InputStream in =
sock.getInputStream();
            BufferedReader bin = new
BufferedReader(new InputStreamReader(in));

            String line;
            while( (line = bin.readLine()) !=
null)

                System.out.println(line);

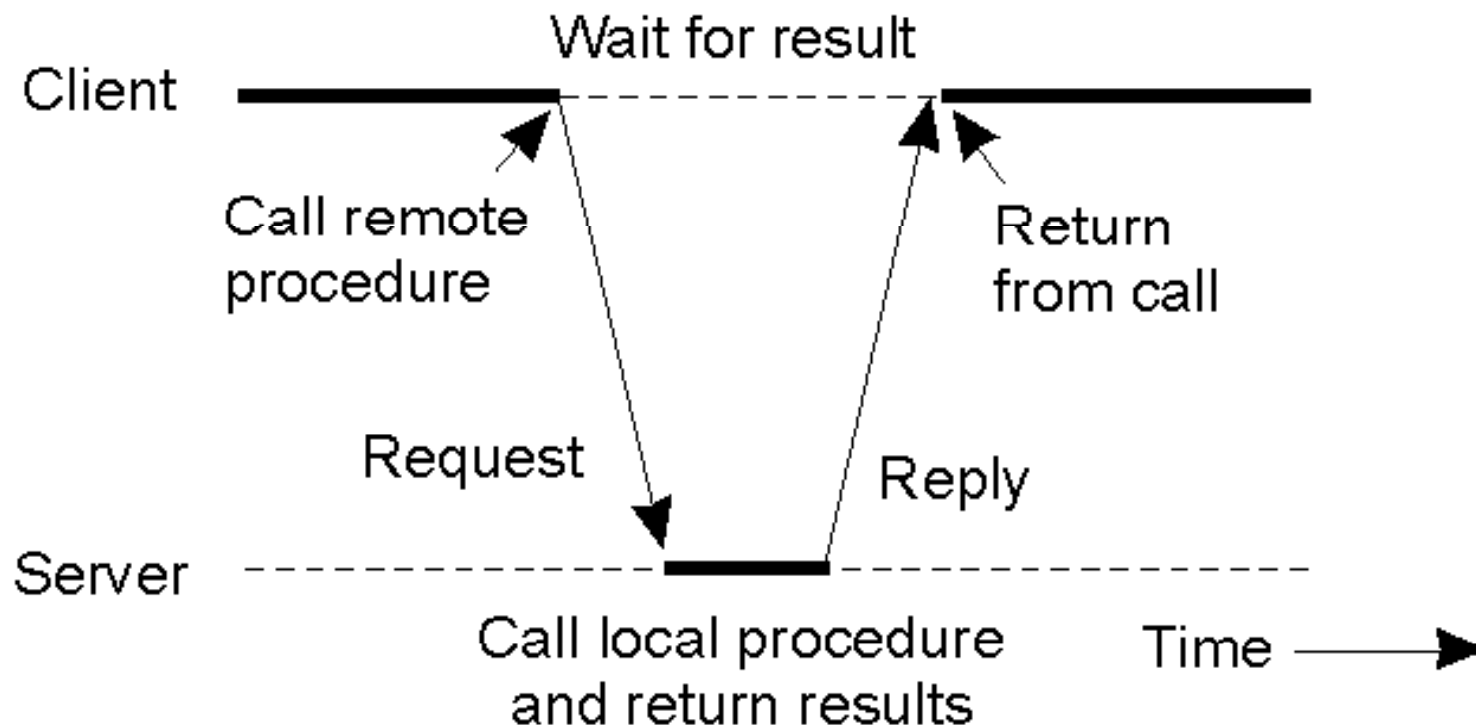
            sock.close();
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

Remote Procedure Call (RPC) (1)

- ❑ Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
- ❑ Stubs
 - Client-side proxy for the actual procedure on the server
- ❑ The client-side stub locates the server and marshalls the parameters
- ❑ The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server

Remote Procedure Call (RPC) (2)

- ❑ RPC between client and server program
- ❑ Request-reply communication



TMS, Fig 2.3

Pipes in Practice

❑ Pipe between two processes

```
ps ax | less or ps ax | grep init or  
ps ax | grep init | awk {'print $1'}
```

❑ Named pipes

```
mkfifo tompipe  
ls -l > tompipe  
ls -l > tompipe
```

❑ Pipes and network with netcat

```
cat web.sh | nc -l 8080  
nc bocek.ch 8080 > web.sh (unencrypted!!)
```