
Chapter 4: Threads

Threads

❑ Objectives

- To introduce the notion of a thread — a fundamental unit of CPU utilization, the basis of multithreaded computer systems
- To briefly discuss the APIs for the Pthreads and Java thread libraries
- To examine issues related to multi-threaded programming

❑ Topics

- Multi-threading Models
- Thread Libraries
- Threading Issues
- Operating System Examples
- Linux Threads

Multi-core Programming

- ❑ All multi-core systems putting pressure on programmers, challenges include
 - Dividing activities
 - Define separate independent concurrent tasks
 - Balance
 - Identify tasks with equal work load
 - Data splitting
 - Divide data to be processed concurrently
 - Data dependency
 - Avoid serial dependencies of computations and synchronizations
 - Testing and debugging
 - Unknown parallel execution path and order

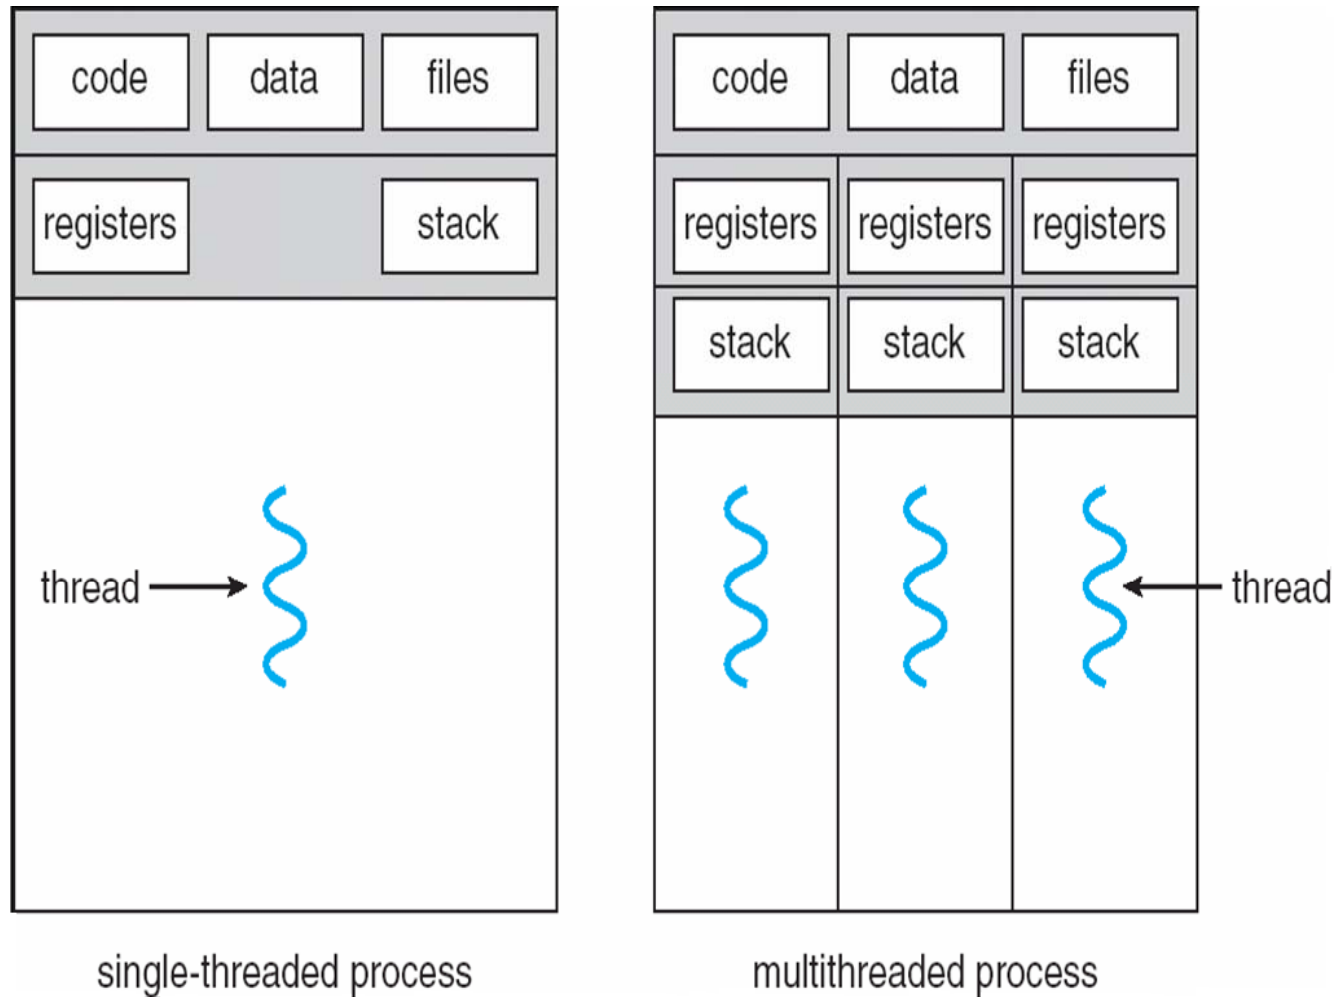
Threads

- ❑ Smallest **unit of CPU utilization** managed by the operating system scheduler
- ❑ Threads are **lightweight processes**
 - Sharing code, data, and files
 - But have separate registers and stacks
- ❑ Threads are contained inside a process
 - Multiple threads can exist within the same process and share resources, while different processes do not share their resources

Comparison: Threads and Processes

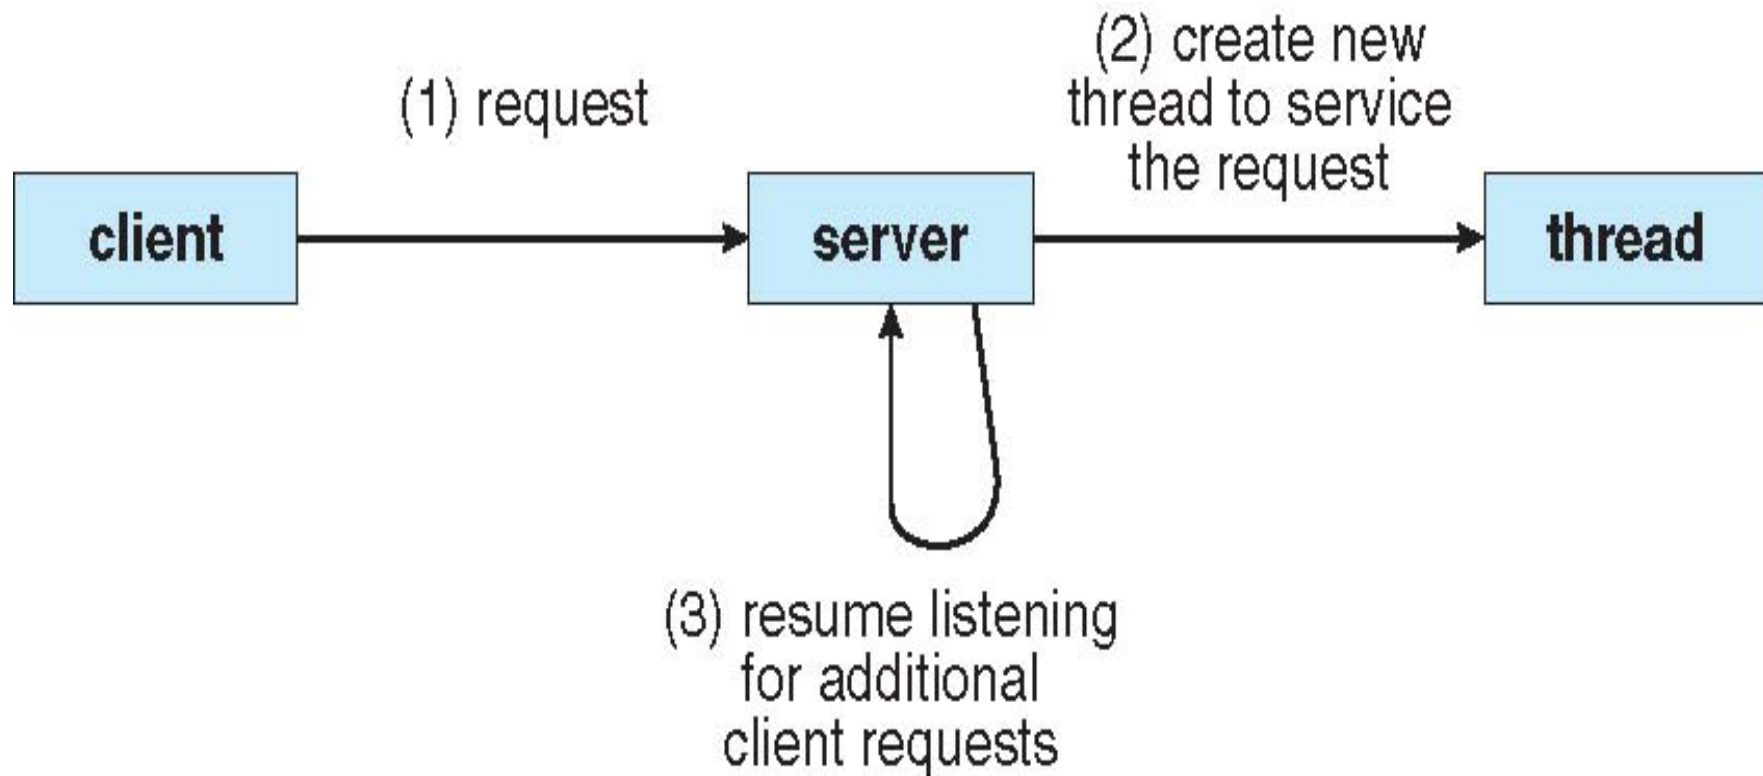
- Processes are typically independent, while threads exist as subsets of a process
- Processes carry considerably more state information than threads, whereas multiple threads within a process share process state as well as memory and other resources
- Processes have separate address spaces, whereas threads share their address space
- Processes interact only through system-provided inter-process communication mechanisms
- Context switching between threads in the same process is typically faster than context switching between processes

Single and Multi-threaded Processes




Benefits: Responsiveness – Resource Sharing – Economy – Scalability

Multi-threaded Server Architecture




Threaded Single/Multi-core Systems

- ❑ Concurrent execution on **single core**



- ❑ Parallel execution on **multi-core (here: 2 cores)**





Challenges: Dividing Activities – Balance – Data Splitting – Data Dependency – Testing and Debugging

Kernel Threads

- ❑ At least **one kernel thread** exists within each process
 - Created and scheduled by the OS kernel
- ❑ **One kernel thread assigned to each (logical) CPU core**
 - Unit of independent process, thread scheduling
- ❑ Support for kernel threads
 - Linux
 - Mac OS X
 - Solaris
 - Tru64 UNIX
 - Since Windows XP/2000

User Threads

- ❑ User and application program thread management done above kernel by **user-level threads library**
- ❑ Ultimately, a relationship must exist between user threads and kernel threads
 - **Many-to-one** prohibits multithreading within user process
 - No concurrent execution as only one user thread can run on a CPU core
 - All user threads get blocked if underlying kernel thread is blocked
 - **One-to-one** and many-to-many implementations enable multi-processor parallelism



Threading (1)

- ❑ Semantics of `fork()` and `exec()` system calls
 - Does `fork()` duplicate only the calling thread or all threads?
- ❑ Thread cancellation of target thread
 - Asynchronous cancellation terminates the target thread immediately
 - Deferred cancellation allows the target thread to periodically check, if it should be cancelled
- ❑ Thread pools
 - Usually slightly faster to service a request with an existing thread than creating a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool

Threading (2)

❑ Thread-specific data

- Allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (*i.e.*, when using a thread pool)

❑ Signal handling

- ... next slide

Signal Handling

- ❑ Signals are used in UNIX systems to notify a process that a particular event has occurred
- ❑ A signal handler is used to process signals
 - Signal is generated by particular event
 - Signal is delivered to a process
 - Signal is handled
- ❑ Handling options
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process

Thread Libraries

- ❑ Thread libraries provide programmers with an API for creating and managing threads
- ❑ Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS
- ❑ Three very common user-level thread libraries
 - POSIX Pthreads
 - Win32 threads
 - Java threads
- ❑ Operating thread systems:
 - Linux

Pthreads

- ❑ May be provided either as user-level or kernel-level
- ❑ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
 - http://standards.ieee.org/findstds/interps/1003-1c-95_int/ (1996)
 - API specifies behavior of thread library
 - Implementation is up to the developer of the library
- ❑ Common in UNIX-like operating systems
 - Solaris, Linux, Mac OS X
 - Win32 threads quite similar to Pthreads

IEEE: Institute of Electrical and Electronics Engineers

Java Threads

- ❑ Java threads are managed by the JVM
 - Threads determine the fundamental program execution model in Java, combined with a rich set of features for creation of threads and their management
- ❑ Typically implemented using the threads model provided by underlying OS
- ❑ Java threads may be created by two alternatives
 - Create a new class derived from the Thread class and override `run()` method
 - Implementing the Runnable interface

Linux Threads

- ❑ Linux refers to them as *tasks* rather than threads
- ❑ Thread creation is done through `clone()` system call
 - Processes are created by `fork()` system call (Chapter 3)
- ❑ `clone()` allows a child task to share the address space of the parent task (process)

flag	meaning
<code>CLONE_FS</code>	File-system information is shared.
<code>CLONE_VM</code>	The same memory space is shared.
<code>CLONE_SIGHAND</code>	Signal handlers are shared.
<code>CLONE_FILES</code>	The set of open files is shared.