

---

# Chapter 5: Scheduling

# CPU Scheduling

---

## ❑ Objectives

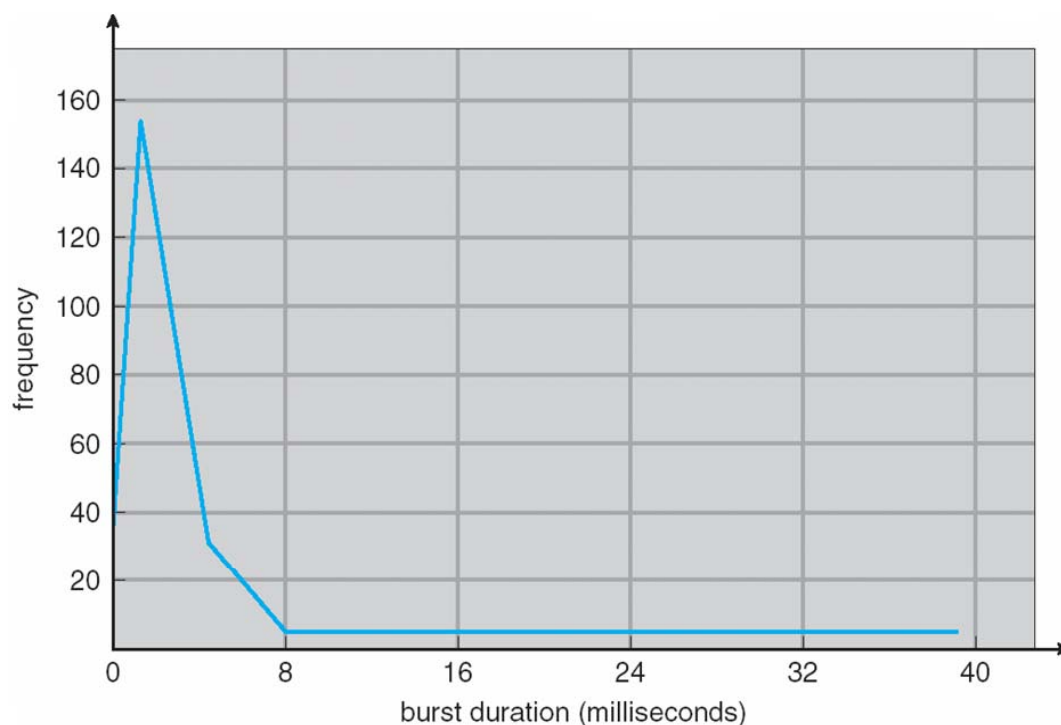
- To introduce CPU scheduling, basis for multi-programmed operating systems
- To describe selected CPU scheduling algorithms
- To discuss evaluation criteria for selecting a CPU scheduling

## ❑ Topics

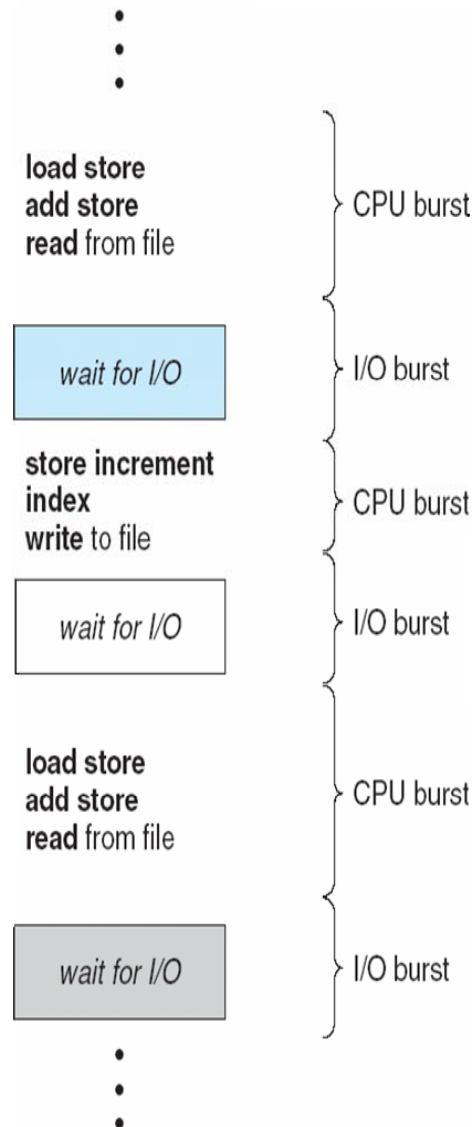
- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Multiple-Processor Scheduling
- Operating Systems Examples

# Basic Concepts

- ❑ Maximum **CPU utilization** only obtained with concurrent multi-programming
- ❑ **CPU-I/O Burst Cycle**
  - Process execution consists of a cycle of CPU execution and I/O wait
- ❑ **CPU burst distribution**
  - Measured extensively
  - Vary greatly from processes, computers, applications

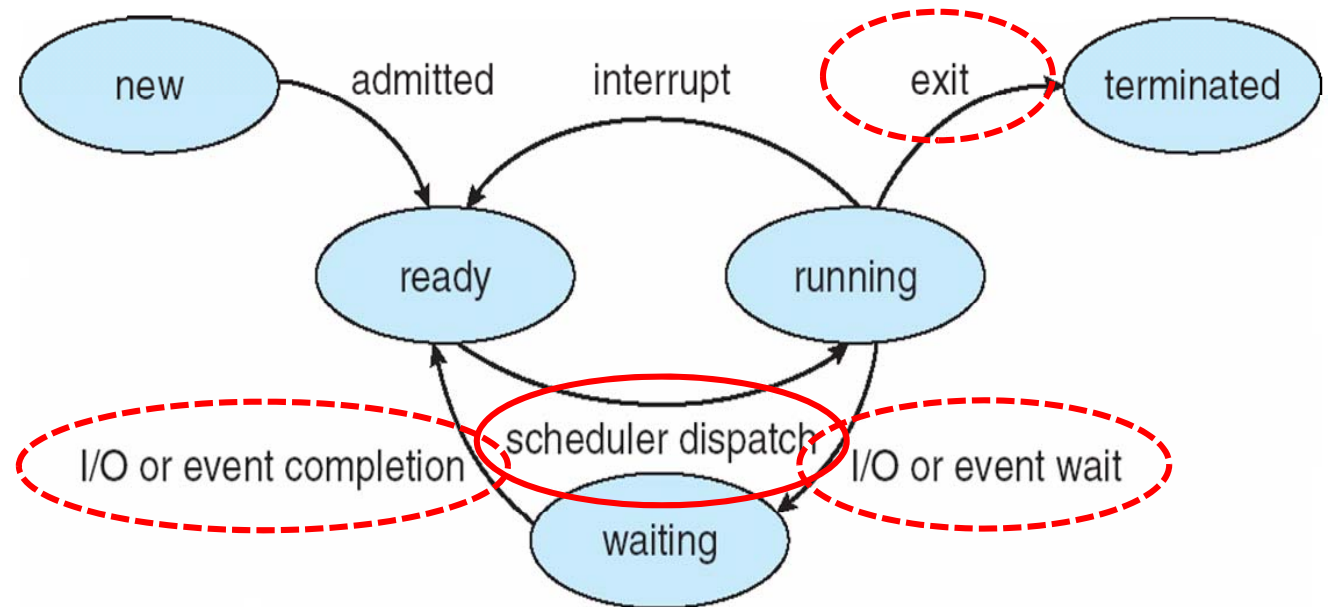


# Alternating Sequence of CPU and I/O Bursts



# Process States (Reminder)

- As a process executes, it frequently changes **states**
  - **new**: the process is being created
  - **running**: instructions are being executed
  - **waiting**: the process is waiting for some event to occur
  - **ready**: the process is waiting to be assigned to a processor
  - **terminated**: the process has finished execution



# CPU Scheduler

---

- ❑ Selects from among processes in memory that are **ready to execute** and **allocates CPU** to one of them
- ❑ Scheduling **decisions** may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state
  3. Switches from waiting to ready
  4. Terminates
- ❑ Scheduling under 1 and 4 is **non-preemptive**
  - Processes cooperate by giving up the CPU voluntarily
- ❑ All other scheduling is **preemptive**
  - Running CPU execution interrupted, triggered by outside event

# Dispatcher Module

---

- ❑ **Dispatcher module** gives control of the CPU to the process selected by the short-term scheduler
- ❑ This involves:
  - Switching context
  - Switching to user mode
  - Jumping to the proper location in the user program to restart that program
- ❑ **Dispatch latency**
  - Time it takes for the dispatcher to stop one process and start another running
  - Must be very low to avoid overhead

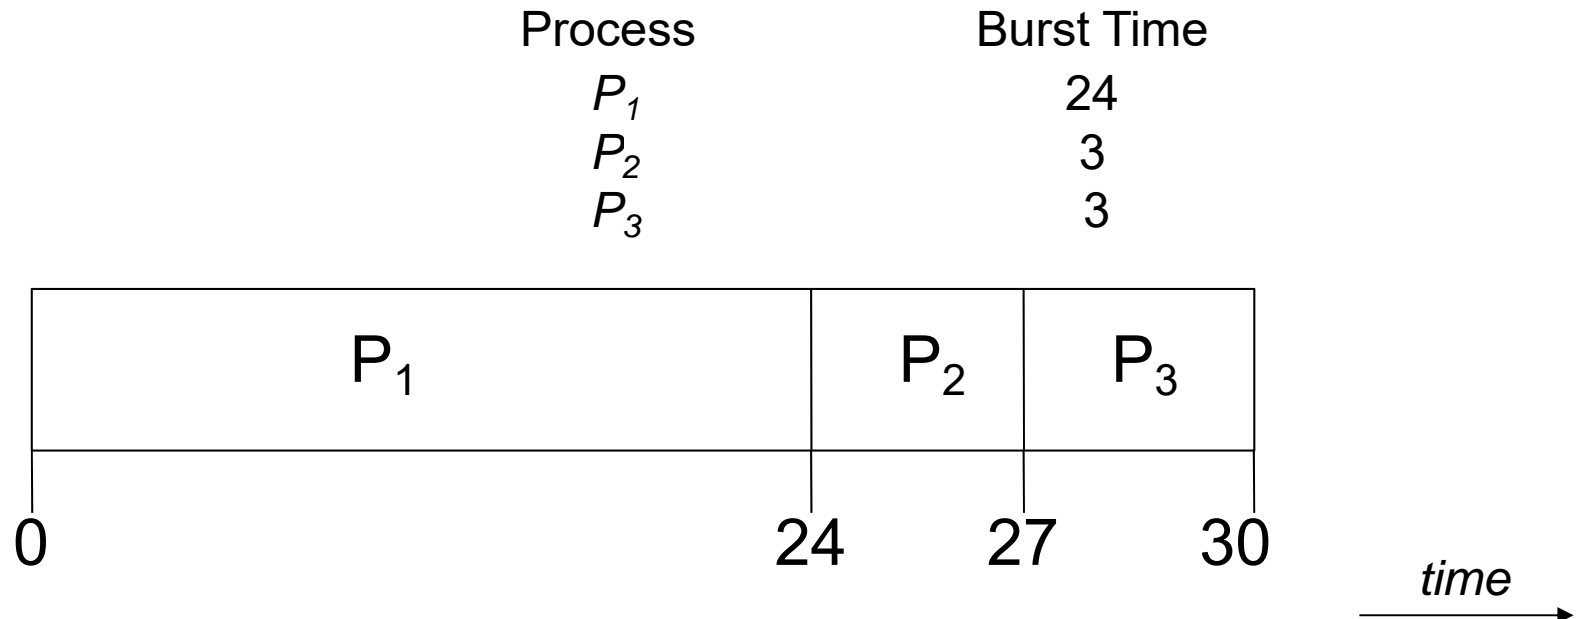
# Scheduling Criteria

---

- ❑ **CPU utilization** – keep the CPU as busy as possible
- ❑ **Throughput** – # of processes that complete their execution per time unit
- ❑ **Turnaround time** – amount of time to execute a particular process
- ❑ **Waiting time** – amount of time a process has been waiting in the ready queue
- ❑ **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)



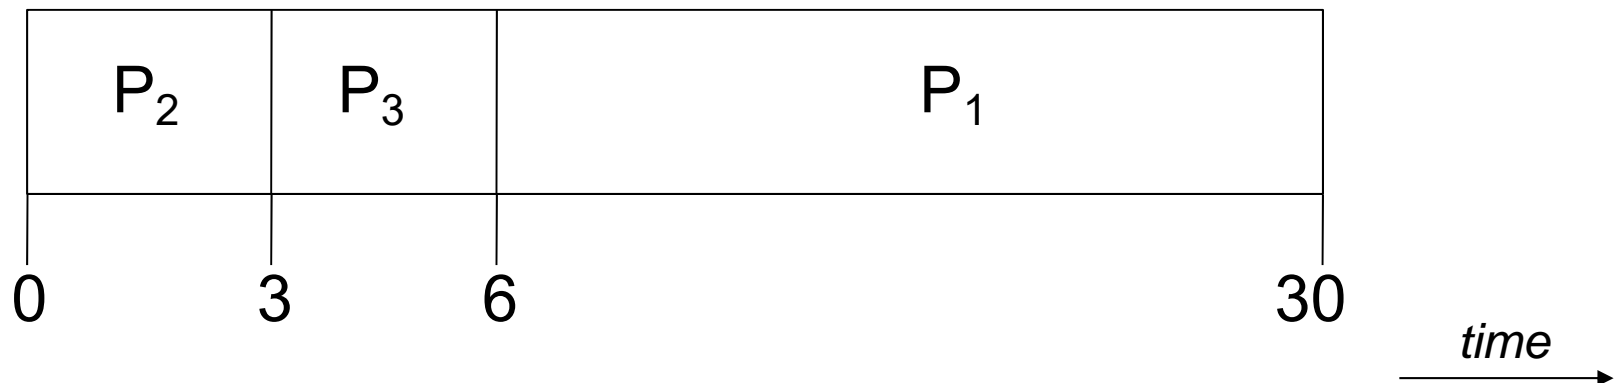
# First Come First Serve (FCFS) Scheduling



- ❑ Suppose that processes arrive in order  $P_1$ ,  $P_2$ ,  $P_3$
- ❑ The **Gantt chart** (bars for start/end times) for the schedule is
  - Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
  - Average waiting time:  $(0 + 24 + 27)/3 = 17$

## FCFS Scheduling (2)

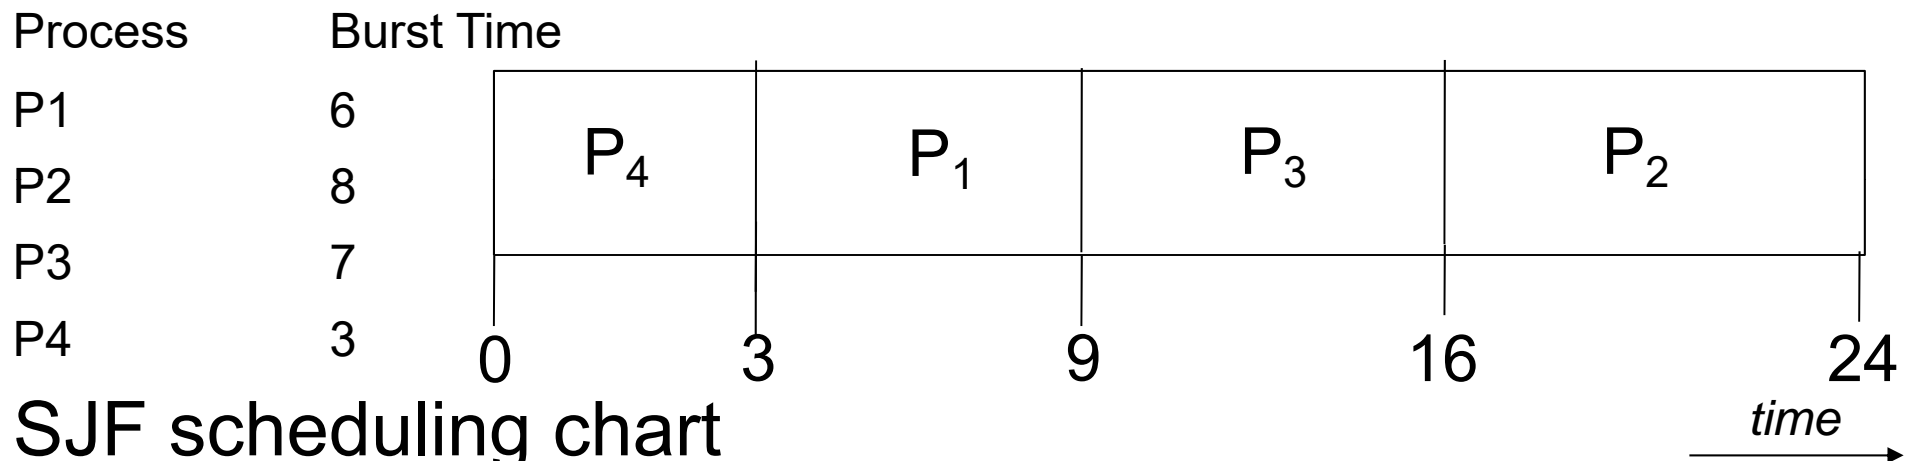
- Suppose that processes arrive in order  $P_2, P_3, P_1$



- The **Gantt chart** for the schedule is
  - Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
  - Average waiting time:  $(6 + 0 + 3)/3 = 3$ 
    - Much better than previous case
  - Convoy effect on short processes behind long processes

# Shortest Job First (SJF) Scheduling

- ❑ Associate with each process the **length of its next CPU burst** and use these lengths to schedule the process with the shortest time to run and complete first
- ❑ SJF is optimal
  - Gives minimum average waiting time for given set of processes
  - The difficulty is knowing the length of the next CPU request



- ❑ SJF scheduling chart

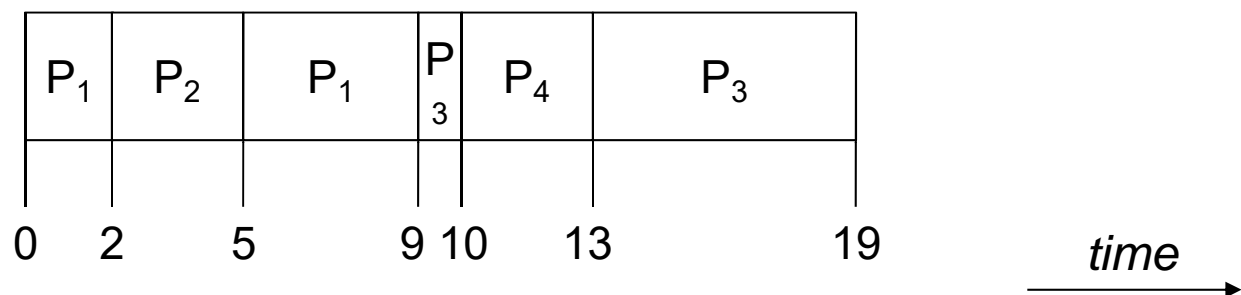
– Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$

# Shortest Remaining Time (SRT) Scheduling

- **Preemptive** SJF scheduling, allowing *new* processes to interrupt the currently running one

Process	Arrival Time	Burst Time
$P_1$	0.0	6
$P_2$	2.0	3
$P_3$	4.0	7
$P_4$	10.0	3

- SRT scheduling chart



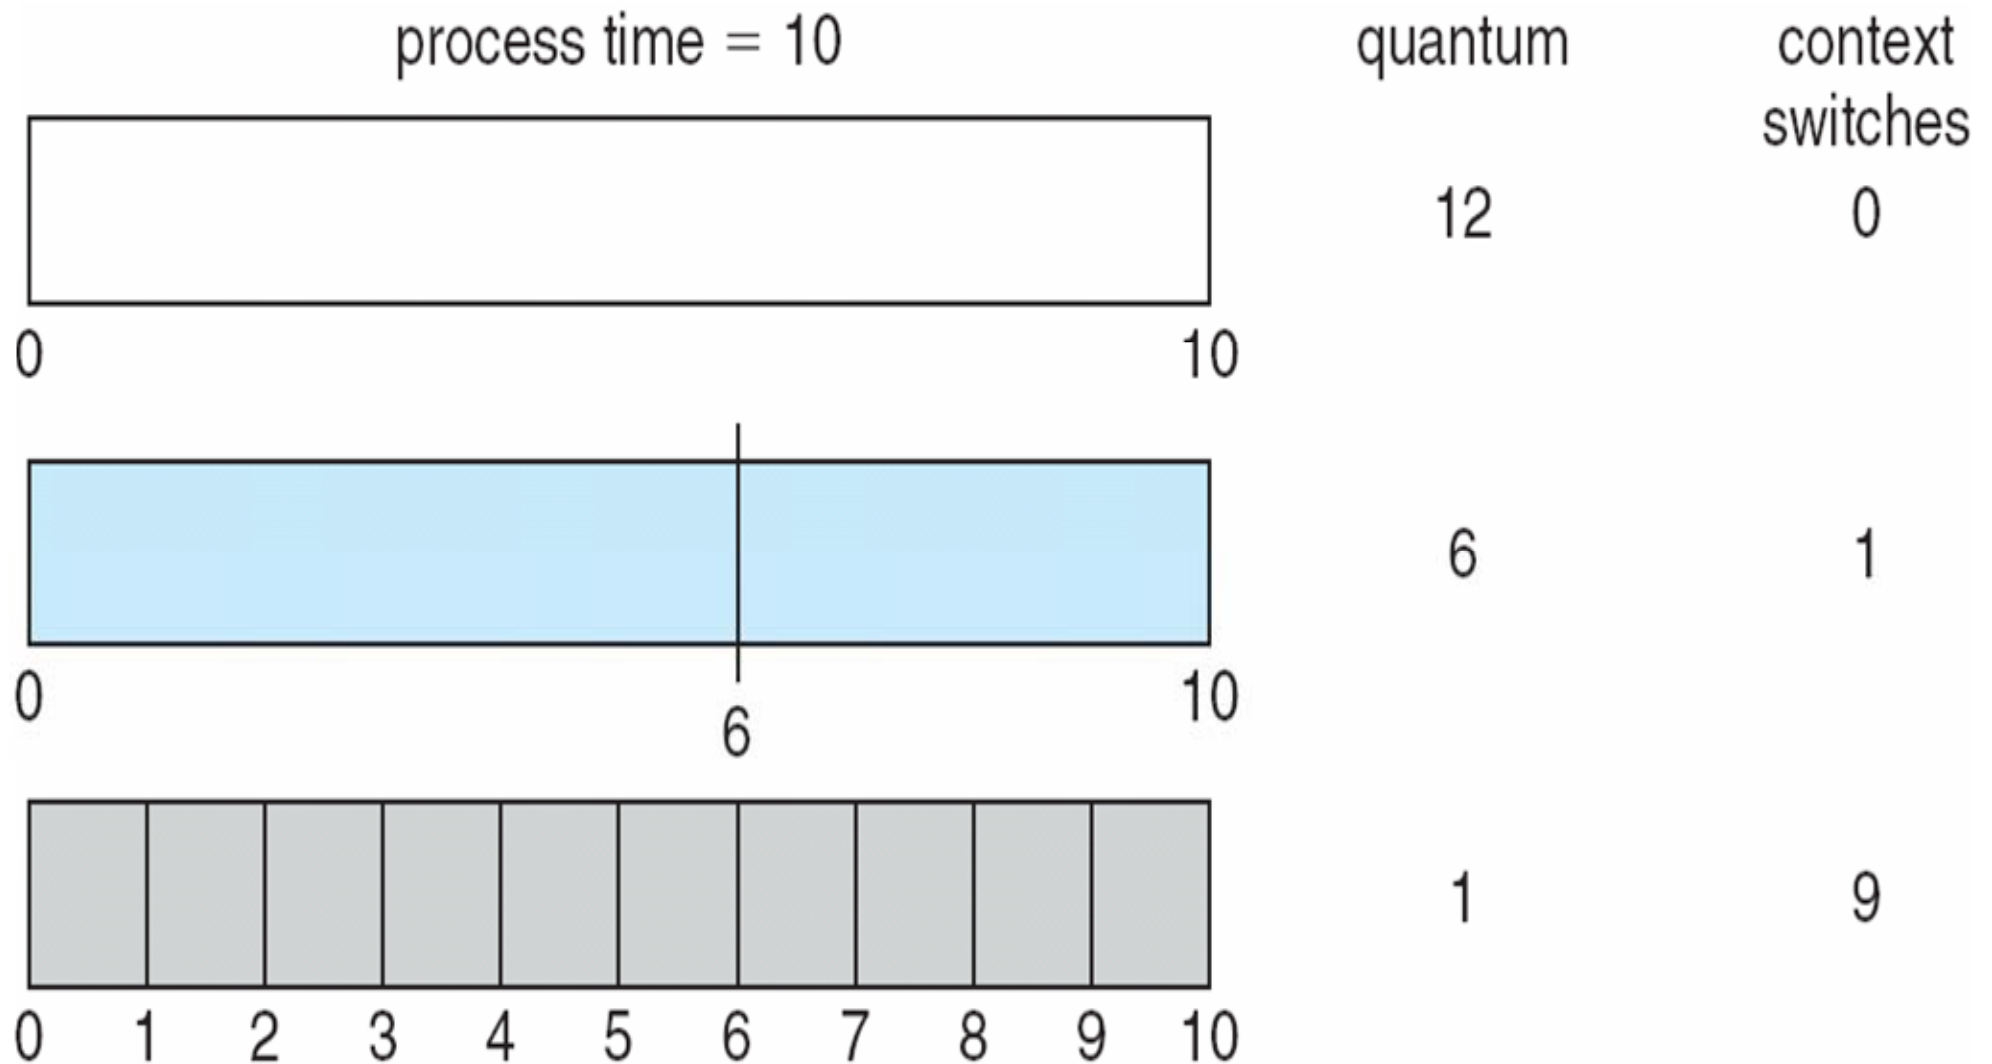
– Average waiting time =  $(3 + 0 + 5 + 3 + 0) / 4 = 2.75$

# Round Robin (RR) Scheduling

---

- ❑ Process gets a small unit of CPU time (*time quantum*),
  - After time elapsed (usually 10-100 ms), process preempted and added to the end of the ready queue
- ❑ If there are  $n$  processes in the ready queue and the time quantum is  $q$ , each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once.
  - No process waits more than  $(n-1)q$  time units
  - $n$  processes run in “parallel” at  $1/n$  of available CPU speed
- ❑ Performance
  - $q$  too large  $\Rightarrow$  FIFO
  - $q$  too small  $\Rightarrow$  many context switches, overhead is too high
    - $q$  must be large with respect to context switch costs

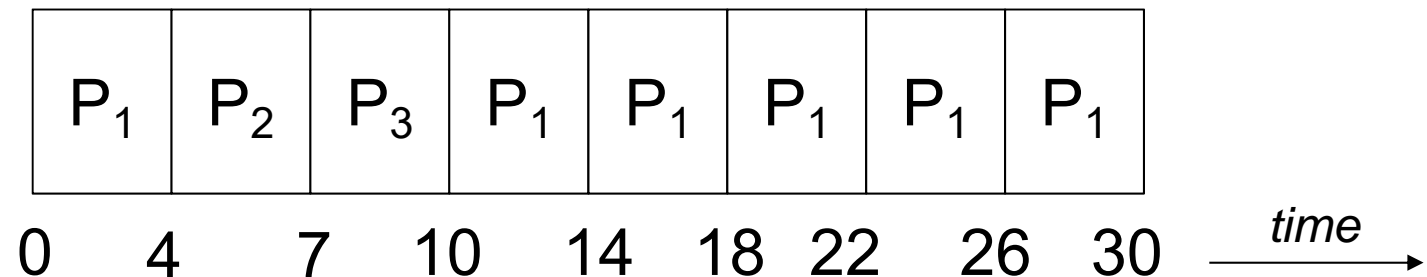
# Time Quantum and Context Switch Time



# Example of RR with Time Quantum = 4

Process	Burst Time
$P_1$	24
$P_2$	3
$P_3$	3

- The Gantt chart is



- Typically, higher average **turnaround** than SJF, but better **response time**
  - Average RR turnaround time  $(30 + 7 + 10) / 3 = 15.6$
  - Comparison to SJF turnaround time  $(3 + 6 + 30) / 3 = 13$

# Priority Scheduling

---

- ❑ Priority number (integer) is associated with each process
- ❑ CPU is allocated to the process with the **highest priority** (smallest integer  $\equiv$  highest priority)
- ❑ FCFS is a priority scheduling, where priority is the process' arrival time: Non-preemptive
- ❑ SJF is a priority scheduling, where priority is the (predicted) next CPU burst time: Preemptive
- ❑ Problem  $\equiv$  **Starvation**
  - Low priority processes may never execute
- ❑ Solution  $\equiv$  Aging
  - As time progresses increase the priority of the process



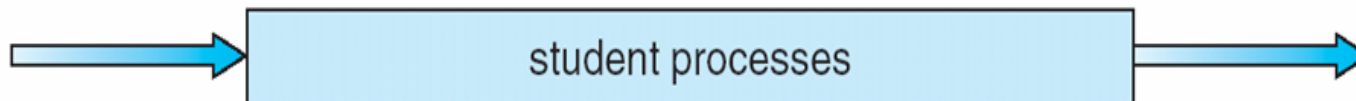
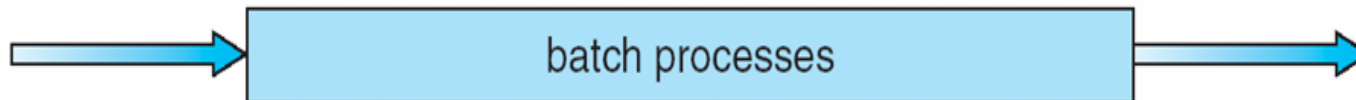
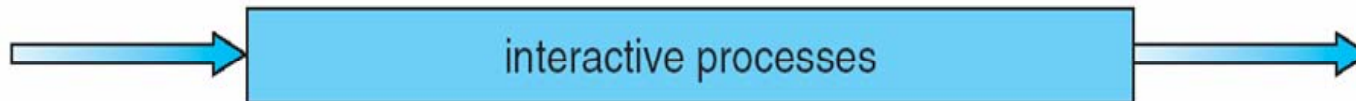
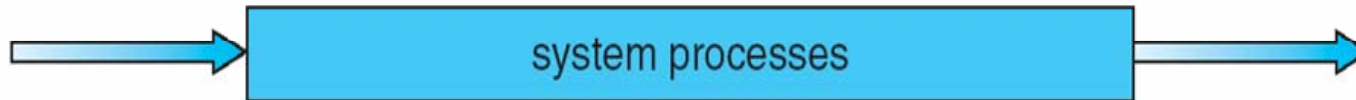
# Multi-level (ML) Queue Scheduling

---

- ❑ Ready queue is partitioned into **separate queues**, *e.g.*:
  - Foreground (interactive)
  - Background (batch)
- ❑ Each queue has its **own scheduling algorithm**
  - Foreground – RR
  - Background – FCFS
- ❑ Scheduling must be done between the queues
  - Fixed priority scheduling, (*i.e.*, serve all from foreground then from background): Possibility of starvation
  - Time slice: each queue gets a certain amount of CPU time, which it can schedule amongst its processes
    - *I.e.*, 80% to foreground in RR and 20% to background in FCFS

# Example for ML Queue Scheduling

highest priority



lowest priority

# Multi-level Feedback Queue (MLF)

---

- ❑ A process can move between the various priority queues
  - Aging can be implemented this way
- ❑ Multi-level feedback queue scheduler defined by the following parameters
  - Number of queues
  - Scheduling algorithms for each queue
  - Method used to determine when to upgrade a process
  - Method used to determine when to demote a process
  - Method used to determine which queue a process will enter when that process needs service

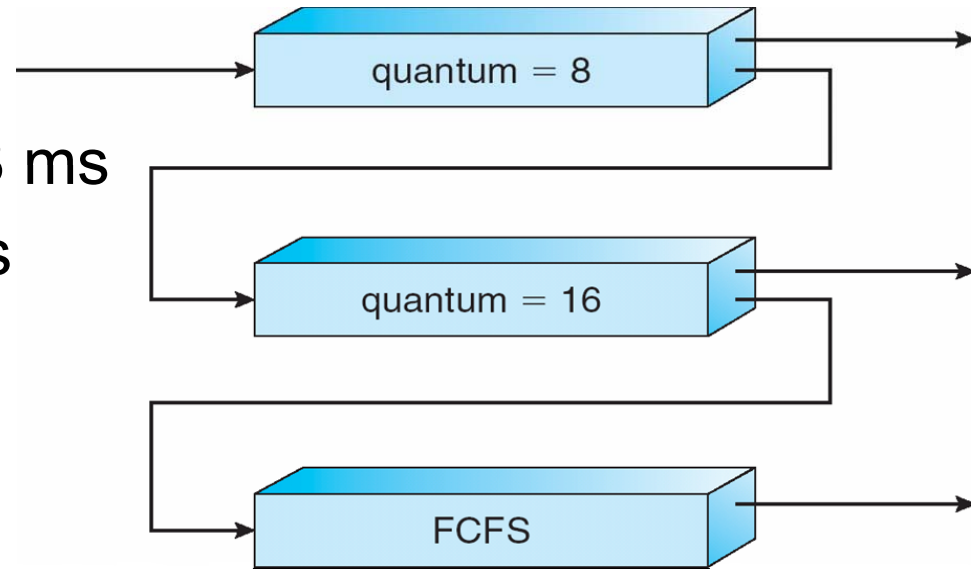
# Example of a 3-Level MLF

## ❑ Three queues

- $Q_0$  – RR with time quantum 8 ms
- $Q_1$  – RR time quantum 16 ms
- $Q_2$  – FCFS

## ❑ Scheduling

- New job enters queue  $Q_0$  which is served RR. When it gains CPU, job receives 8 ms. If it does not finish in 8ms, job is preempted and moved to queue  $Q_1$
- At  $Q_1$  job is again served RR and receives 16 additional ms. If it still does not complete, it is preempted and moved to queue  $Q_2$
- Final queue  $Q_2$  is served FCFS



# Multiple Processor Scheduling (1)

---

- ❑ CPU scheduling is more complex when multiple CPUs are available
  - Homogeneous processors within a multi-processor assumed
- ❑ Asymmetric multiprocessing (AMP)
  - Only one processor accesses system data structures
  - Alleviating the need for data sharing
- ❑ Symmetric multiprocessing (SMP)
  - Each processor is self-scheduling
  - All processes in common ready queue or each has its own private queue of ready processes
  - Supported, e.g., in Windows XP, Mac OS X, Linux, Solaris

# Multiple Processor Scheduling (2)

---

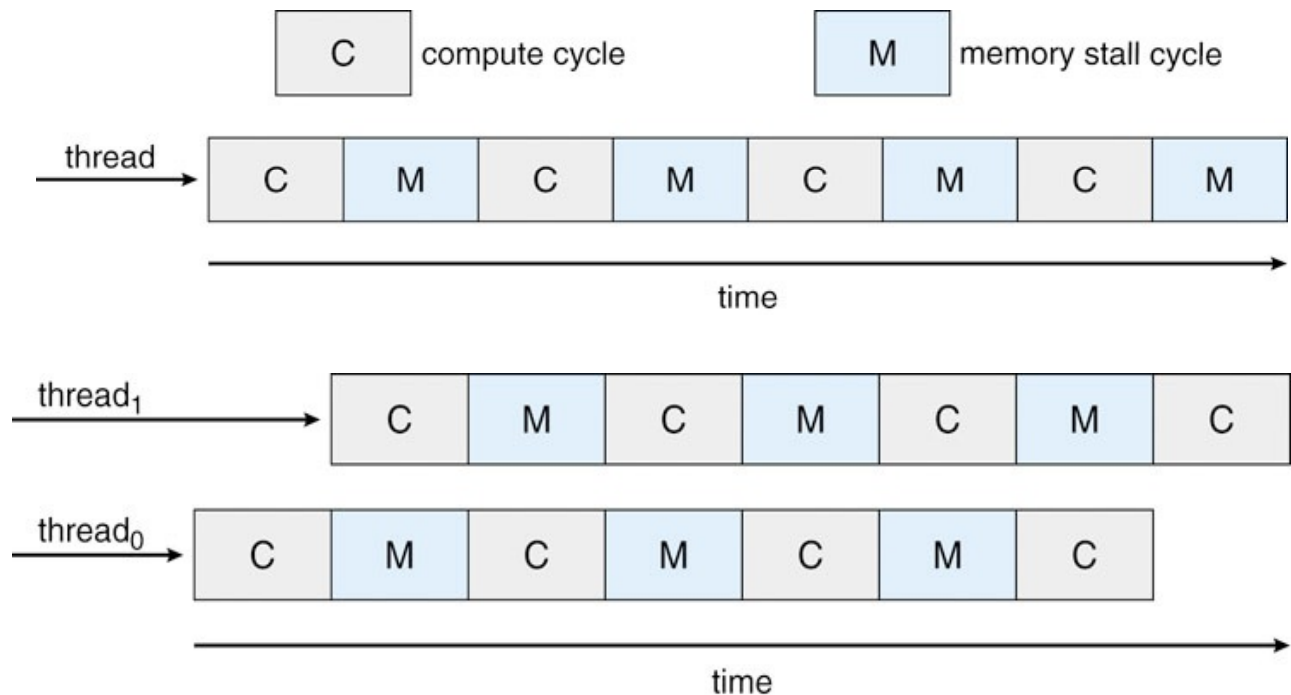
## ❑ Processor affinity

- Process has affinity for processor on which it is currently running
  - Hard affinity
    - Process runs only on one particular CPU
  - Soft affinity
    - Process can migrate to another CPU, but reluctantly done

## ❑ Process migration involves cache invalidation on one CPU and repopulation on the other

# Multi-core Processors

- ❑ Trend since 2010 to place multiple processor cores on same physical chip
  - Faster, consumes less power
- ❑ Multiple kernel threads per core also growing
  - Takes advantage of memory stall to make progress on another thread, while memory retrieve happens



# Example 1: Linux Scheduling/Priorities (1)

---

- ❑ Constant order  $O(1)$  scheduling time independent of number of processes
  - Suitable for large and SMP systems
  
- ❑ Two priority ranges
  - Real-time
    - Real-time range from 0 to 99
  - Time sharing (nice values)
    - Nice value from 100 to 140
  - Longer time quantum for higher priority



# Example 1: Linux Scheduling/Priorities (2)

- Time sharing tasks have **dynamic priorities**, adjusted after expired time quantum

- Level of interactivity determines  $\pm 5$  adjustment

	<u>numeric priority</u>	<u>relative priority</u>		<u>time quantum</u>
– Long sleep times				
→ more interactive	0	highest	real-time tasks	200 ms
	•			
	•			
– Short sleep times	99		other tasks	
	100			
→ less interactive	•			
• CPU bound	•			
	140	lowest		10 ms

# Example 2: Java Thread Scheduling

---

- ❑ JVM uses a preemptive, priority-based scheduling algorithm
  - FIFO queue is used if there are multiple threads with the same priority
- ❑ JVM schedules a thread to run when
  - The currently running thread exits the runnable state
  - A higher priority thread enters the runnable state
- ❑ Note: the JVM does not specify whether threads are time-sliced or not