
Chapter 9: Virtual Memory

Virtual Memory

❑ Objectives

- To describe the **benefits** of a virtual memory system
- To explain the **concepts** of demand paging, page-replacement algorithms, and allocation of page frames
- To discuss the **principles** of the working-set model

❑ Topics

- Background
- Demand Paging
- Process Creation: Copy-on-Write and Memory-mapped Files
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Other Considerations

Background


❑ Virtual memory

- Separation of user logical memory from physical memory
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation


❑ Virtual memory can be implemented via

- Demand paging
- Demand segmentation

Virtual Memory and Address Space

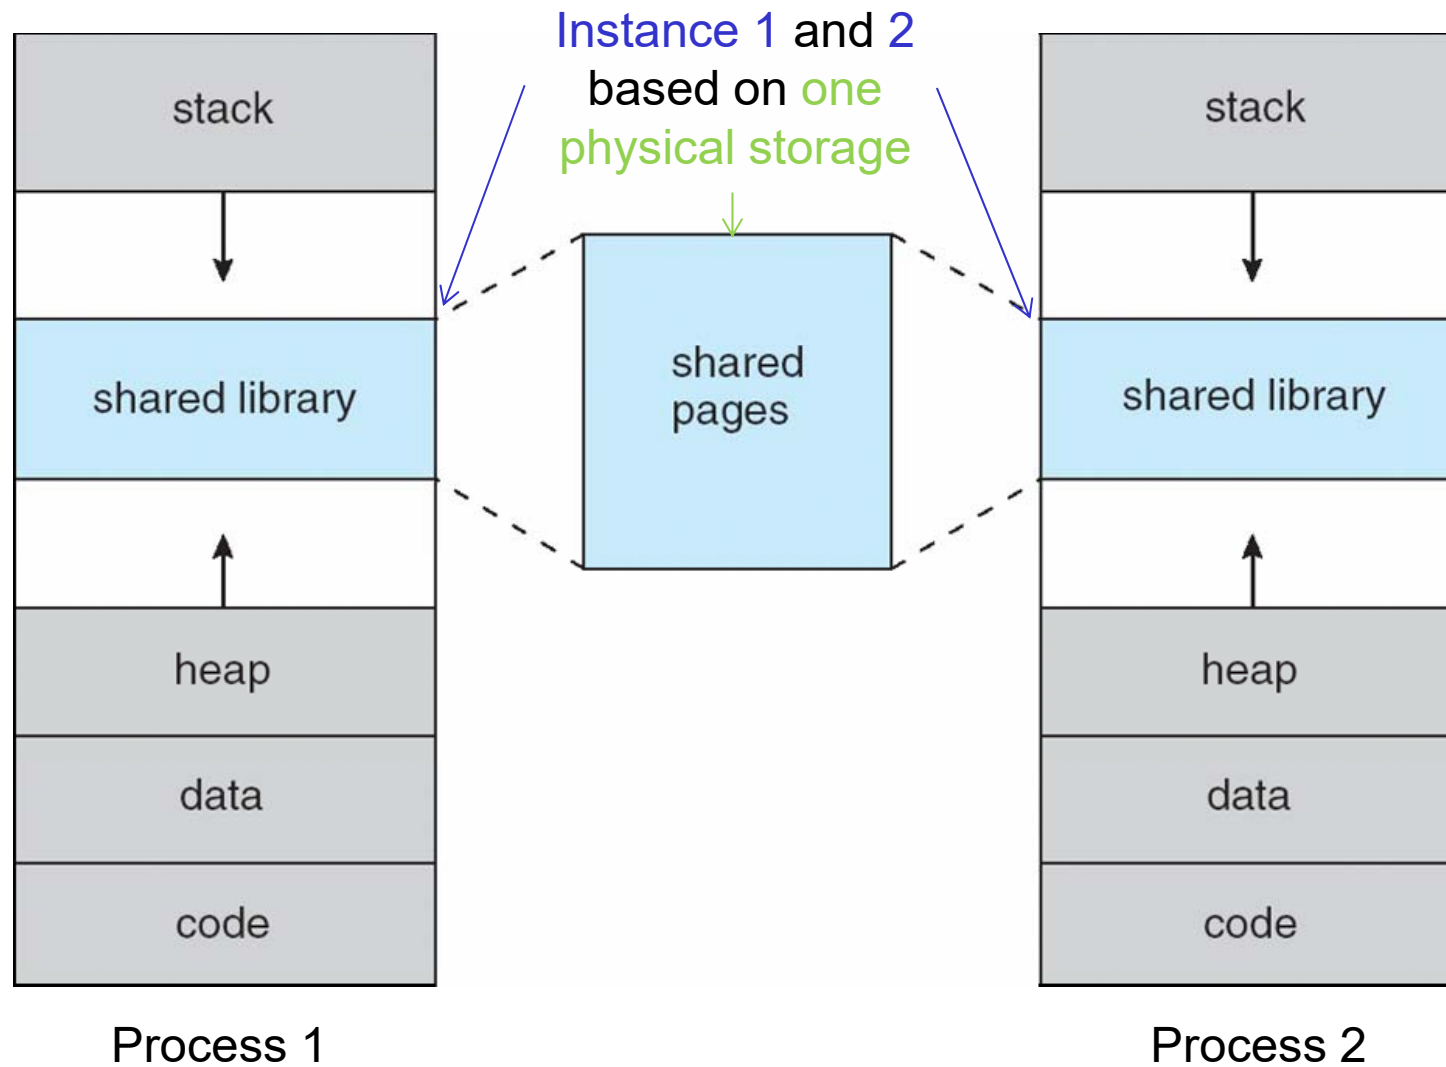


Virtual Memory (larger than Physical Memory)



Virtual Address Space

Shared Library Using Virtual Memory



Demand Paging

- ❑ Bring a **page into memory** only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- ❑ Page is needed \Rightarrow reference to it
 - Invalid reference \Rightarrow abort
 - Not-in-memory \Rightarrow bring to memory
- ❑ **Lazy swapper**: never swaps a page into memory unless page will be needed
 - Swapper that deals with pages is a **pager**

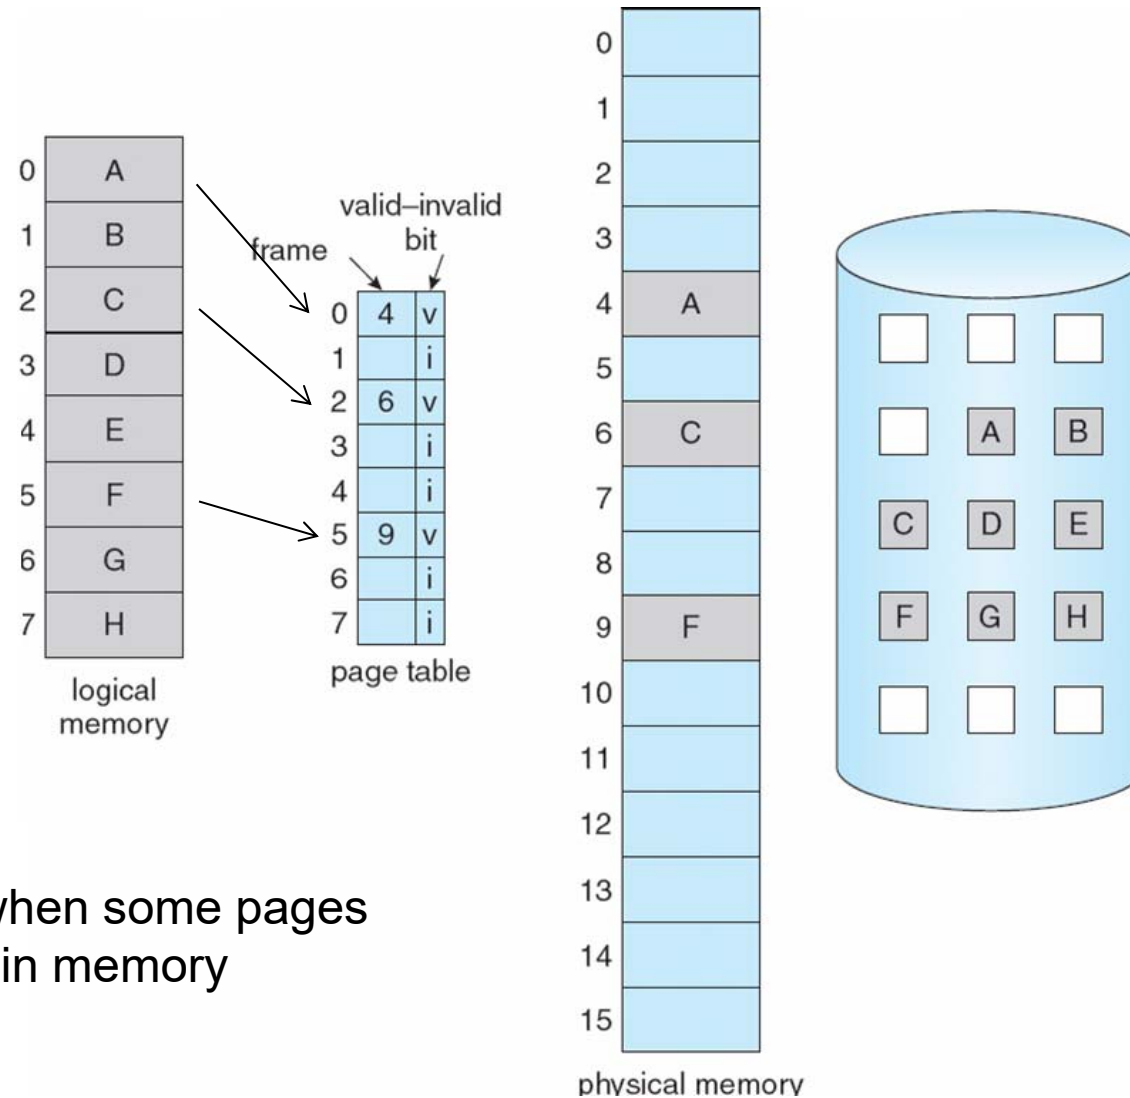
Valid-Invalid Bit

- ❑ With each page table entry a **valid–invalid bit** is associated
 - **v** \Rightarrow in-memory
 - **i** \Rightarrow not-in-memory
- ❑ Initially valid–invalid bit is set to **i** on all entries
- ❑ Example of a **page table snapshot**:
During address translation,
if valid–invalid bit in page table entry is **i** \Rightarrow page fault

Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

Pages Missing in Main Memory

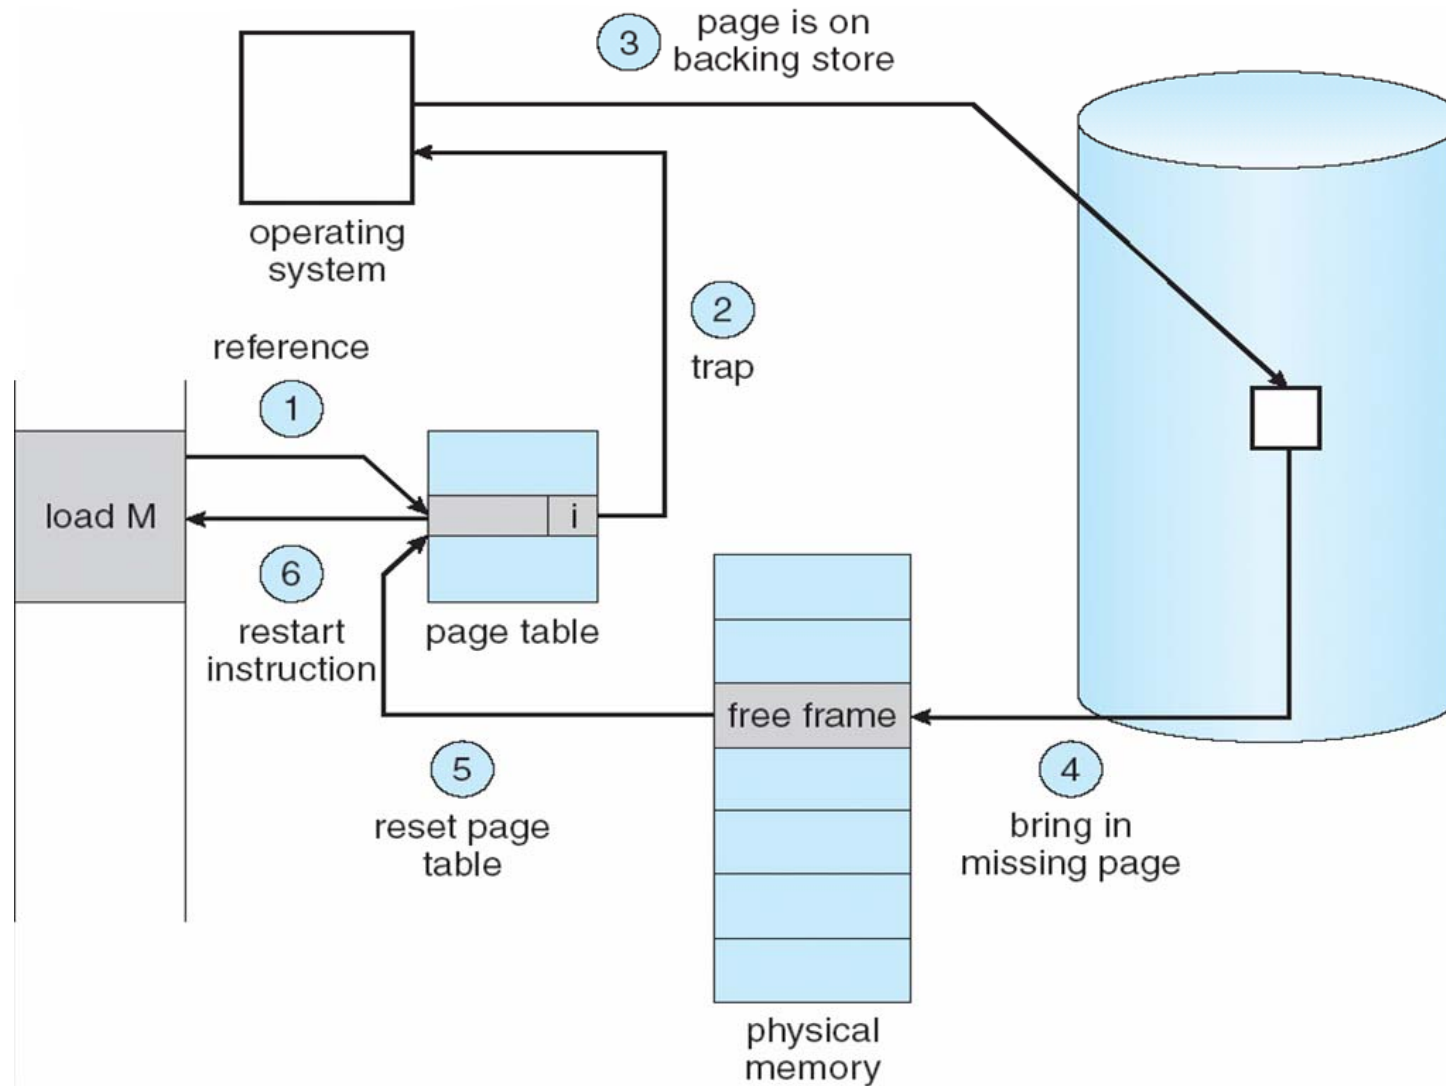


Page Table when some pages are not in main memory

Page Fault

- ❑ If there is a reference to a page, first reference to that page will trap to operating system
 - Page fault
- 1. Operating system looks at another table to decide:
 - Invalid reference \Rightarrow abort
 - Just not in memory
- 2. Get empty frame
- 3. Swap page into frame
- 4. Reset tables
- 5. Set validation bit = **v**
- 6. Restart the instruction that caused the page fault

Steps in Handling a Page Fault



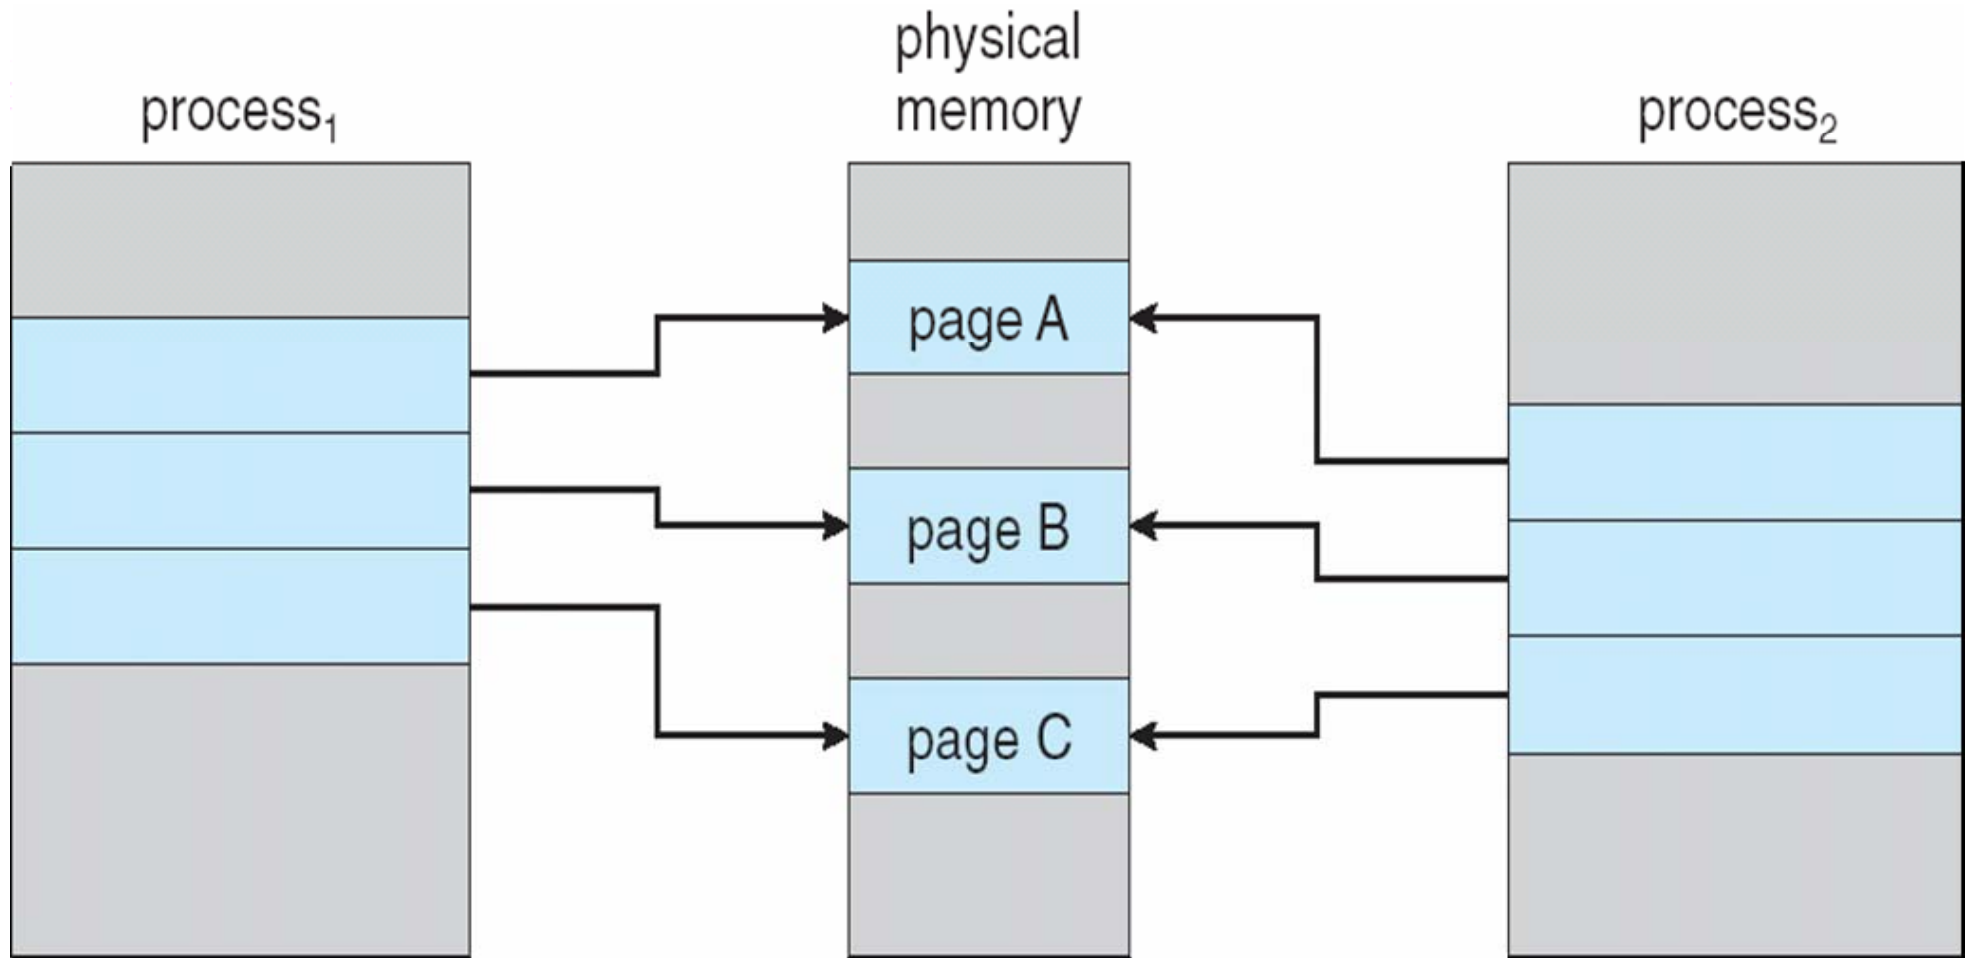
Process Creation

- ❑ Virtual memory allows for other benefits during the process creation (such as `fork()`):
 - Copy-on-Write
 - Memory-Mapped Files

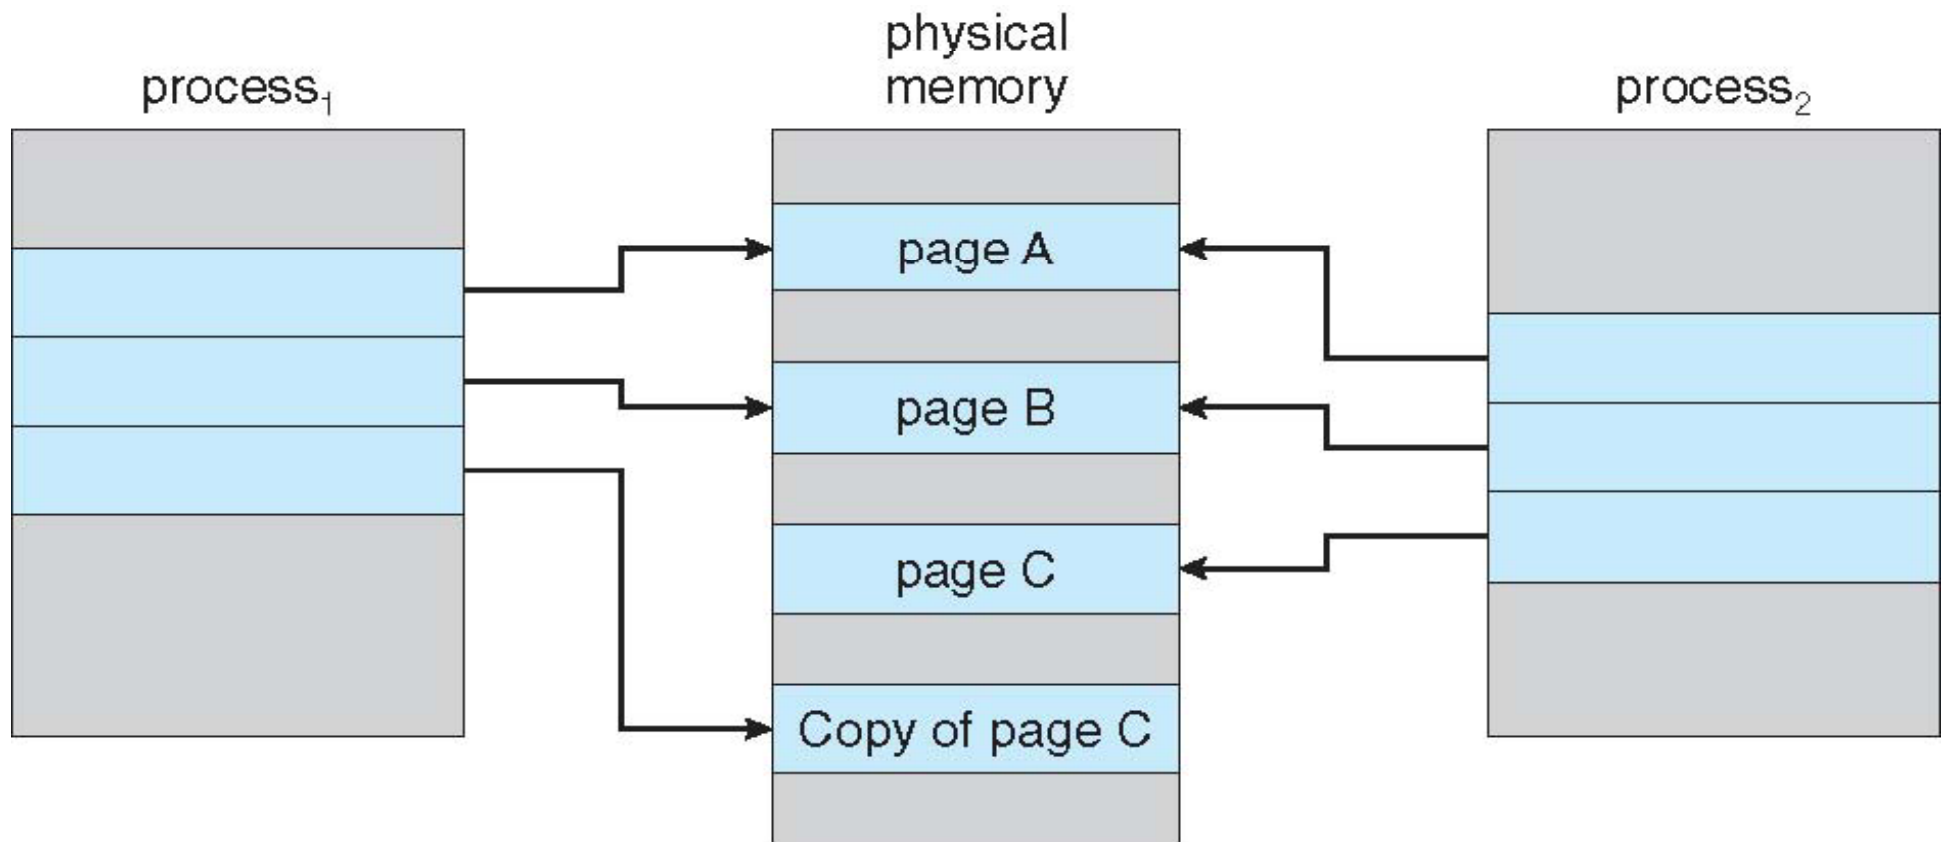
Copy-on-Write

- ❑ Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory
 - If either process modifies a shared page, only then is the page copied
- ❑ COW allows more efficient process creation as only modified pages are copied
- ❑ Free pages are allocated from a **pool** of zeroed-out pages

Before Process 1 Modifies Page C



After Process 1 Modifies Page C



Only pages that can be modified need be marked as copy-on-write!

Lack of Free Frames

❑ Page replacement

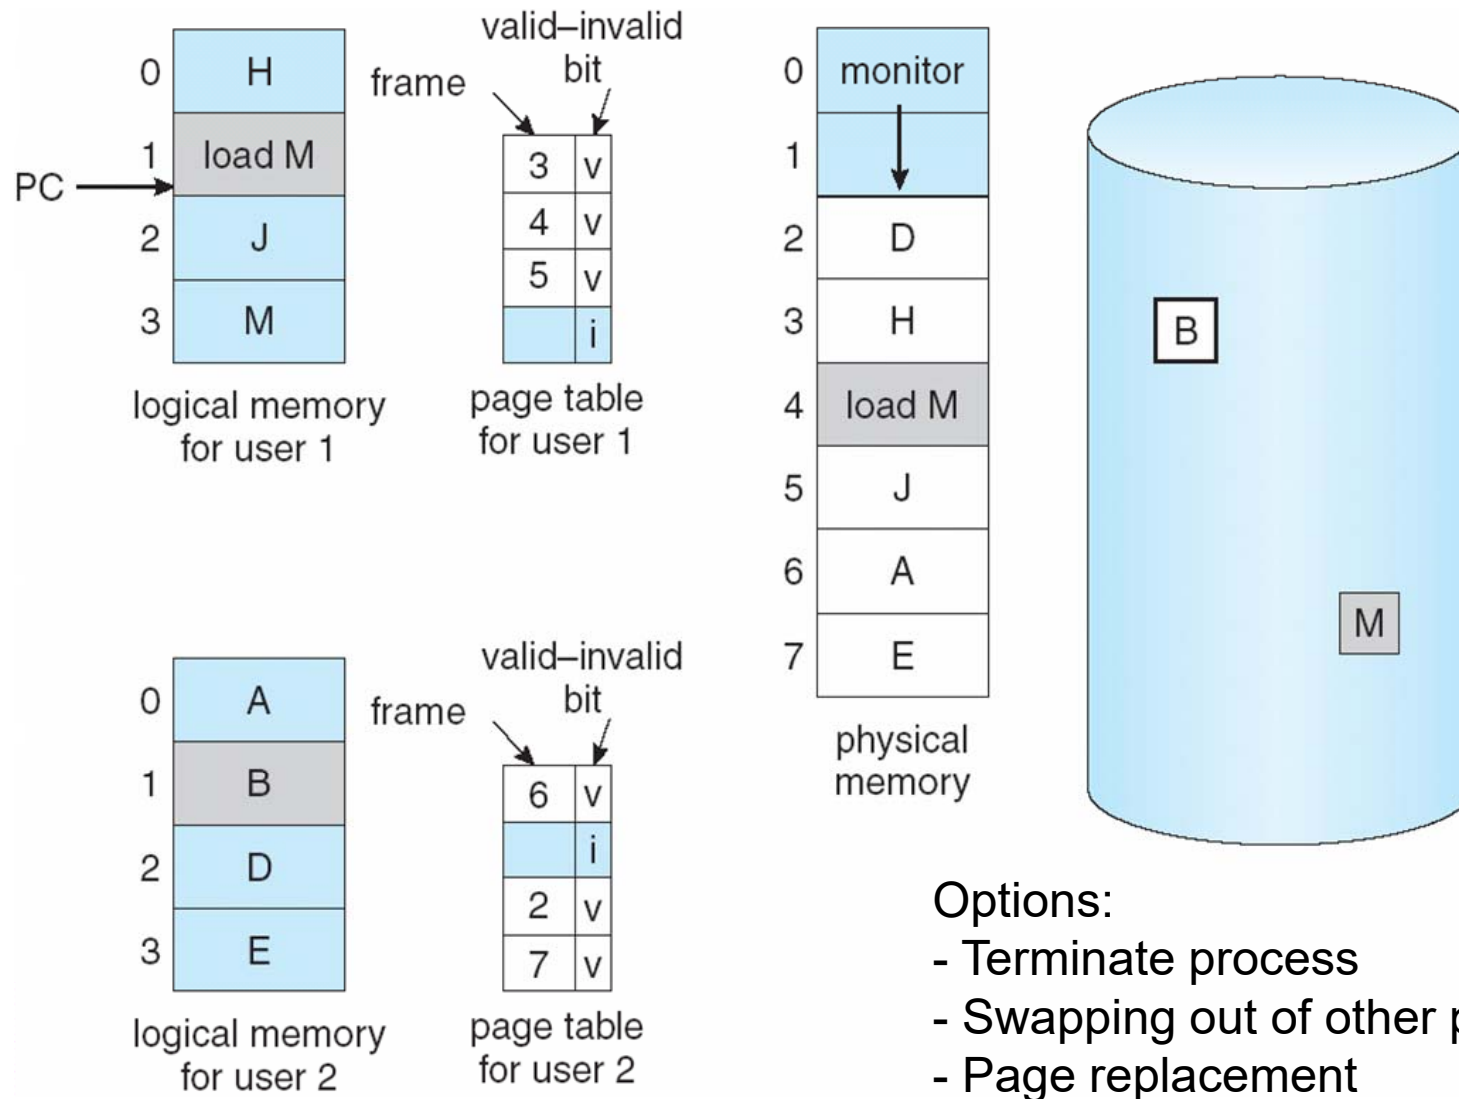
- Find some page in memory, but not really in use, swap it out
 - Algorithm
 - Performance demands
 - In search of an algorithm which will result in a minimum number of page faults

❑ Same page may be brought into memory several times

Page Replacement

- ❑ Prevent over-allocation of memory by modifying page-fault service routine to include page replacement
- ❑ Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- ❑ Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory


Need for Page Replacement



Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Restart the process

Page Replacement




Page Replacement Algorithms

- ❑ Want lowest page-fault rate
- ❑ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
- ❑ In all of the examples the reference string is:

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Page Faults vs. Number of Frames



Expectation: #frames increases, #page faults drops, or adding physical memory

First-in-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process):

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults

4 frames:

1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults

FIFO queue:

- All page requests queued
- Replace page at head of queue
- Page brought into memory is added at tail of queue

- Belady's Anomaly: more frames \Rightarrow more page faults

FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																
	0	0	0																
		1	1																

2	2	4	4	4	0														
3	3	3	2	2	2														
1	0	0	0	3	3														

0	0																		
1	1																		
3	2																		

7	7	7																	
1	0	0																	
2	2	1																	

page frames

Optimal Algorithm

- ❑ Replace page that will not be used for longest period of time
- ❑ 4 frames example:

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	4
2	
3	
4	5


6 page faults

- ❑ How do you know this?
- ❑ Used for measuring how well your algorithm performs

Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

Least Recently Used (LRU) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

- Counter implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - When a page needs to be changed, look at the counters to determine which are to change

LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1	
	0	0	0		0		0	0	3	3			3		0		0	
		1	1		3		3	2	2	2			2		2		7	

page frames

Counting Algorithms

- ❑ Keep a counter of the number of references that have been made to each page
- ❑ LFU Algorithm
 - Replaces page with smallest count
- ❑ MFU Algorithm
 - Based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Allocation of Frames

- ❑ Each process needs *minimum* number of pages
 - How to allocate the fixed amount of free memory among P_i ?
- ❑ Example: IBM 370
 - 6 pages to handle SS MOVE instruction:
 - Instruction is 6 byte, might span 2 pages
 - 2 pages to handle *from*
 - 2 pages to handle *to*
- ❑ Two **major allocation schemes**
 - Fixed allocation
 - Priority allocation

Fixed Allocation

□ Equal allocation

- *E.g.*, if there are 100 frames and 5 processes, give each process 20 frames

□ Proportional allocation

- Allocate according to the size of process

s_i = size of process p_i

$S = \sum s_i$

m = total number of frames

a_i = allocation for $p_i = \frac{s_i}{S} \times m$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

Priority Allocation

- ❑ Use a proportional allocation scheme using priorities rather than size
- ❑ If process P_i generates a page fault
 - Select for replacement one of its frames
 - Select for replacement a frame from a process with lower priority number

Global vs. Local Allocation

❑ Global replacement


- Process selects a replacement frame from the set of all frames; one process can take a frame from another

❑ Local replacement

- Each process selects from only its own set of allocated frames


Thrashing

- ❑ If a process does not have “enough” pages, the page-fault rate is very high
- ❑ This leads to
 - Low CPU utilization
 - Operating system thinks that it needs to increase the degree of multiprogramming
 - Another process added to the system
- ❑ **Thrashing** \equiv a process is busy swapping pages in and out



Page-Fault Frequency Scheme


- ❑ Establish “acceptable” page-fault rate
 - If actual rate too low, process loses frame
 - If actual rate too high, process gains frame



Memory Mapped Files (1)

- ❑ Memory mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory
- ❑ A file is initially read using demand paging; a page-sized portion of the file is read from the file system into a physical page; subsequent reads/writes to/from the file are treated as ordinary memory accesses
- ❑ Simplifies file access by treating file I/O through memory rather than **read() / write()** system calls
- ❑ Also allows several processes to map the same file allowing the pages in memory to be shared

Memory Mapped Files (2)



Other Issues (1)

□ Prepaging

- To reduce the large number of page faults that occurs at process startup
- Prepage all or some of the pages a process will need, before they are referenced
- But if prepaged pages are unused, I/O and memory was wasted
- Assume s pages are prepaged and α of the pages is used
 - Is cost of $s * \alpha$ save pages faults $>$ or $<$ than the cost of prepping $s * (1 - \alpha)$ unnecessary pages?
 - α near zero \Rightarrow prepping loses

Other Issues (2)

- ❑ Page size selection must take into consideration:
 - fragmentation
 - table size
 - I/O overhead
 - Locality

- ❑ I/O Interlock
 - Pages must sometimes be locked into memory
 - Consider I/O: Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm