
Chapter 14: Coordination

Coordination

❑ Objectives


- To develop an overview on atomicity, a major factor in DS
- To overview coordination middleware systems as examples

❑ Topics

- Coordination and dependencies
- Atomicity and failures
- One-Phase Commit (1PC)
- Two-Phase Commit (2PC)
- Middleware Systems
 - Publish-subscribe paradigm
 - TIB/Rendezvous
 - Jini
 - Zookeeper

Coordination and Dependencies

- ❑ The act of coordinating, making different people or things work together for a goal or effect
- ❑ Three types



Key:  **Resource**  **Activity**

- Co-location
- Daily build
- Interface definitions

- Push-based
- Pull-based
- Market

- First-come first-served
- Manager decides
- Bidding

Atomicity

- ❑ Consider a replicated database
 - Running on a distributed system with reliable multicasting
 - Updates are multicast to all replicas and the system guarantees that they are delivered in order
- ❑ Nasty scenario:
 - A message is multicast reliably to all replicas and is delivered to the application layer (the database)
 - One replica crashes, while performing the update
 - When it recovers it is in an inconsistent state!
- ❑ **Atomicity** is the key and needed
 - All commit or all abort!
 - **Guarantee** that an operation is completed at all participants or at none

Is this sufficient?

Example

- ❑ Transfer money from bank A to bank B
 - Debit A, credit B, tell client “OK”
- ❑ This process wants **either both** to perform the transfer or **neither** to do it
 - Never want only one side to act!
 - Better if nothing happens!
- ❑ Goal: **Atomic Commit Protocol**

Two Major Aspects of Atomicity

❑ Serializability

- Series of operations requested by users
- Outside observer sees each of them complete atomically in some complete order
- Requires support for locking

❑ Recoverability

- Each operation executes completely or not at all
- No partial results exist – “all-or-nothing”

❑ Prerequisites

- Serializability applies synchronization
 - Logical and vector clocks (see earlier module)
- Recoverability applies a distributed protocol (to follow)

Difficulty of Atomic Commits

❑ Example


- A tells B: “I’ll commit, if you commit”
- A hears no reply from B
- Now what?
- Neither party can make final decision!

❑ Impacts from

- Communication systems errors
- Communication protocols’ unreliability
- Distributed systems’ hardware failures
- Application misbehaviors

One-Phase Commit (1PC) Protocol

- ❑ Create a Transaction Coordinator (TC)
 - A single authoritative entity
- ❑ Four entities
 - Client, TC, Bank A, Bank B
- ❑ Operation
 - Client sends “start” to TC
 - TC sends “debit” to A
 - TC sends “credit” to B
 - TC reports “OK” to client



Failure Scenarios

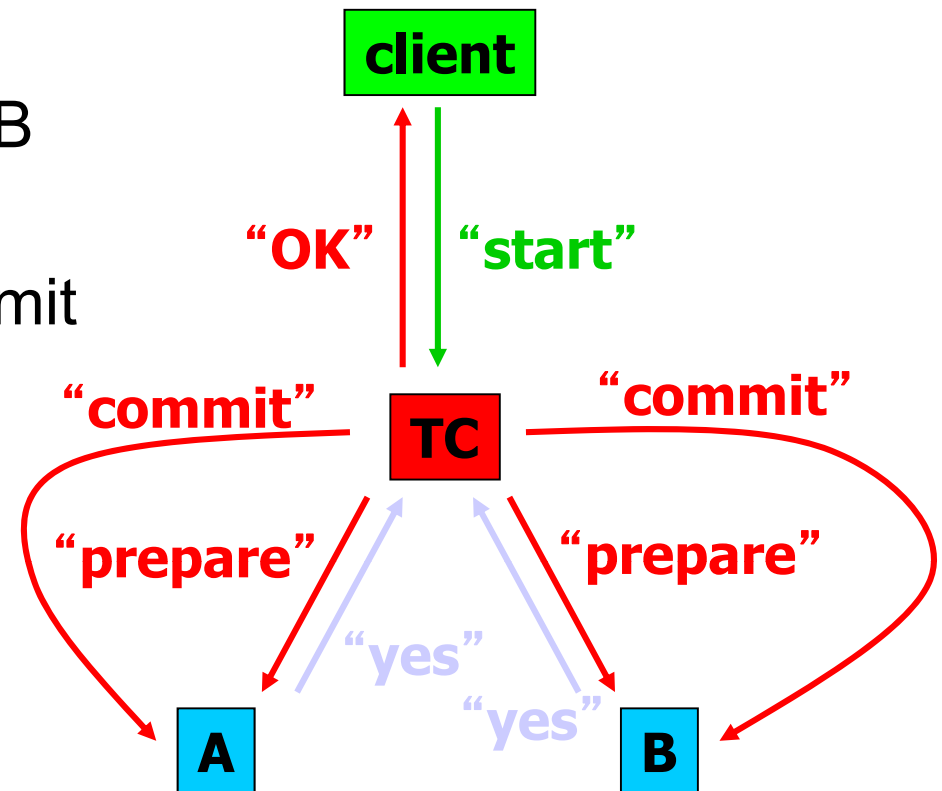
- ❑ Not enough money in A's bank account
 - A does not commit, B does
- ❑ B's bank account no longer exists
 - A commits, B does not
- ❑ One network link (of A or B) is broken
 - One commits, the other does not
- ❑ One of A or B has crashed
 - One commits, the other does not
- ❑ TC crashes between sending to A and B
 - A commits, B does not

Desirable Properties of an Atomic Commit

- ❑ TC, A, and B have separate notions of committing
- ❑ Correctness
 - If one commits, no one aborts
 - If one aborts, no one commits
- ❑ Liveness (in a sense related to performance)
 - If no failures and A and B can commit, then commit
 - If failures, come to “some” conclusion as soon as possible

Two-Phase Commit (2PC) Protocol

- ❑ Same entities as in 1PC
- ❑ Operation
 - Client “starts” and TC sends “prepare” messages to A and B
 - A and B respond, saying whether they’re willing to commit
 - If both say “yes,” TC sends “commit” messages
 - If either says “no,” TC sends “abort” messages
 - A and B “decide to commit”, if they receive a commit message.



Correctness and Liveness of 2PC

- ❑ Why is the 2PC protocol correct (*i.e.*, safe)?
 - Knowledge **centralized at TC** about willingness of A and B to commit
 - TC enforces both and must agree for either to commit

- ❑ Does the previous protocol always complete (*i.e.*, does it exhibit **liveness**)?
 - No!
 - What if nodes crash or messages get lost?

Liveness Problems in 2PC

□ Timeout

- Host is up, but does not receive message it expects
- Maybe other host crashed, maybe network dropped message, maybe network is down
- Usually cannot distinguish these cases, so solution must be correct for all!

□ Reboot

- Host crashes, reboots, and must “clean up”
- *i.e.*, want to wind up in correct state despite reboot

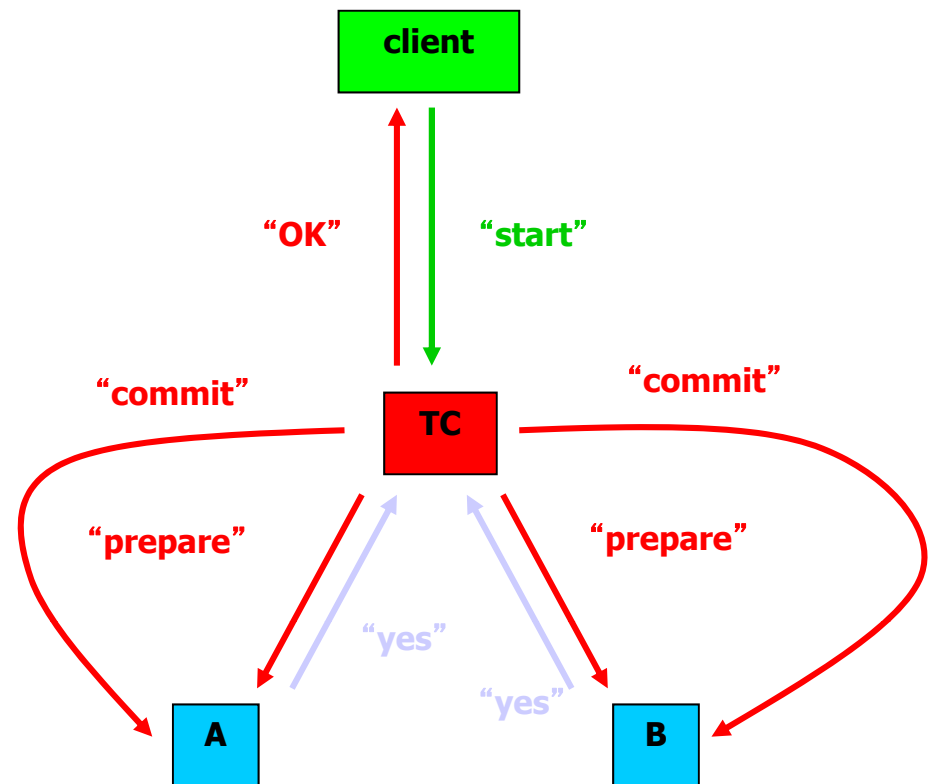
Solution to Liveness Problems

□ Solution

- Introduce timeouts
- Take appropriate actions
 - But be conservative to preserve correctness!

□ Where in the protocol do **hosts wait** for messages?

- TC waits for “yes”/”no” from A and B
- A and B wait for “commit”/”abort” from TC



Transaction Coordinator Times Outs

- ❑ Proceed with a decision when TC waits too long for *yes/no*
- ❑ TC has not yet sent any “commit” messages
 - So it can safely abort
 - Send “abort” messages
- ❑ This preserves safety, but sacrifices liveness
 - Perhaps both A, B prepared to commit, but a “yes” message was lost
 - Could have committed, but TC unaware!
 - Thus, TC is conservative

A (or B) Times Out

- ❑ If B voted “no”
 - It can unilaterally abort
 - TC will never send “commit” in this case
- ❑ If B voted “yes”
 - B cannot unilaterally abort
 - TC might have received “yes” from both, sends “commit” to A, then crashed before sending “commit” to B
 - Result: A would commit, B would abort: incorrect (unsafe)!
 - B cannot unilaterally commit
 - A might have voted “no”
- ❑ If B voted “yes”
 - Either B keeps waiting forever (not a solution) or ...

Better plan?

Termination Protocol (B with “yes” Vote)

- ❑ B directly contacts A: sends “status” request to A asking, if A knows whether the transaction should commit
 - If A received “commit” or “abort” from TC:
 - B decides same way (cannot disagree with TC)
 - If A hasn’t voted anything yet: B and A both abort
 - TC can’t have decided “commit”; it will eventually hear from A or B
 - If A voted “no”: B and A both abort
 - TC can’t have decided “commit”
 - If A voted “yes”: no decision possible, keep waiting
 - TC might have timed out, aborted, and replied to client
 - If no reply from A: no decision possible, wait for TC

Termination Protocol Behavior

- ❑ Some timeouts can be resolved with a guaranteed correctness (safety)
- ❑ Sometimes, though, A and B must block
 - Especially when TC fails or TC's network connection fails
- ❑ Remember
 - TC is entity with **centralized knowledge** of A's and B's state!

Problem: Crash-and-Reboot

- ❑ Cannot back out if commit once decided
 - Suppose TC crashes just after deciding and sending a “commit”
 - What if “commit” message to A or B is lost?
 - Suppose A and/or B crash just after sending “yes”
 - What if “yes” message to TC is lost?

- ❑ If A or B reboots, they do not remember saying “yes”, which leads to big trouble!
 - Might change mind after reboot
 - Even after everyone reboots, may not be able to decide!

Solution: Persistent State (1)

- ❑ Storing **state in non-volatile memory** (e.g., a disk)
 - If all nodes know their pre-crash state, they can use the previously described termination protocol
 - A and B can also ask TC, which may still remember if it committed

- ❑ The order of store and send
 - Write disk
 - Send “yes” message if A/B or “commit” if TC?
 - Or vice-versa?

Solution: Persistent State (2)

- ❑ Can a message be send before writing the disk?
 - Might then reboot between sending and writing, and change mind after reboot
 - *E.g.*, B might send “yes”, then reboots, then decides “no”

- ❑ Thus, write disk before sending message?
 - For TC, write “commit” to disk before sending
 - For A/B, write “yes” to disk before sending

Revised Recovery Protocol

- ❑ TC: after reboot, if no “commit” on disk, abort
 - No “commit” on disk means no “commit” messages had been sent: safe
- ❑ A/B: after reboot, if no “yes” on disk, abort
 - No “yes” on disk means that no “yes” messages had been sent, so no one could have committed; safe
- ❑ A/B: after reboot, if “yes” on disk, use ordinary termination protocol
 - Might block!
- ❑ If everyone rebooted and is reachable, can still decide!
 - Just look at whether TC has stored a “commit” on disk

Summary of 2PC Properties

- ❑ “Prepare” and “commit” phases
 - Two-Phase Commit (2PC)
- ❑ Properties:
 - Safety: All hosts that decide reach the same decision
 - Safety: No commit unless everyone says “yes”
 - Liveness:
 - If no failures occur and all say “yes,” then commit
 - If failures occur, repair, wait long enough, eventually take a decision
- ❑ Remember: **Consensus not (always) possible!**
 - Theorem [Fischer, Lynch, Paterson, 1985]: “No distributed asynchronous protocol can correctly agree (provide both safety and liveness) in presence of crash-failures (*i.e.*, if failures are not repaired)”

Time and Reference Coupling


□ Views of a distributed system

	Time-coupled	Time-uncoupled
Named/ Reference	Communications directed toward a defined receiver that must exist at the same time <i>Examples: RPC, RMI</i>	Communications directed toward a defined receiver that exists at some time <i>Examples: Message based systems</i> A
Unnamed	Unknown name of receiver who must exist at the same time <i>Examples: Multicast</i> B	Unknown name of receiver who should exist at some time <i>Examples: Tuple Spaces, Zookeeper</i> C

A Network of Brokers

A

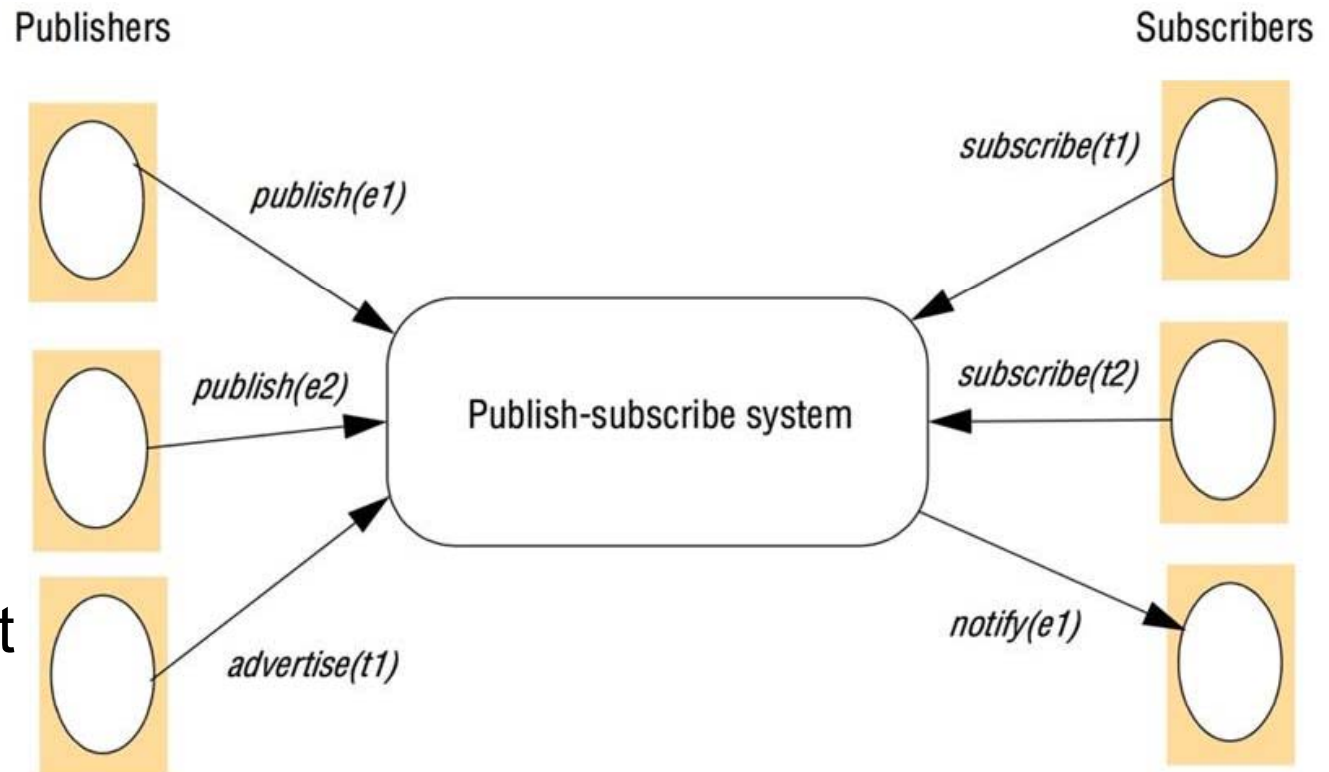
- ❑ Message-based systems
 - Publishers and subscribers interconnected by broker network



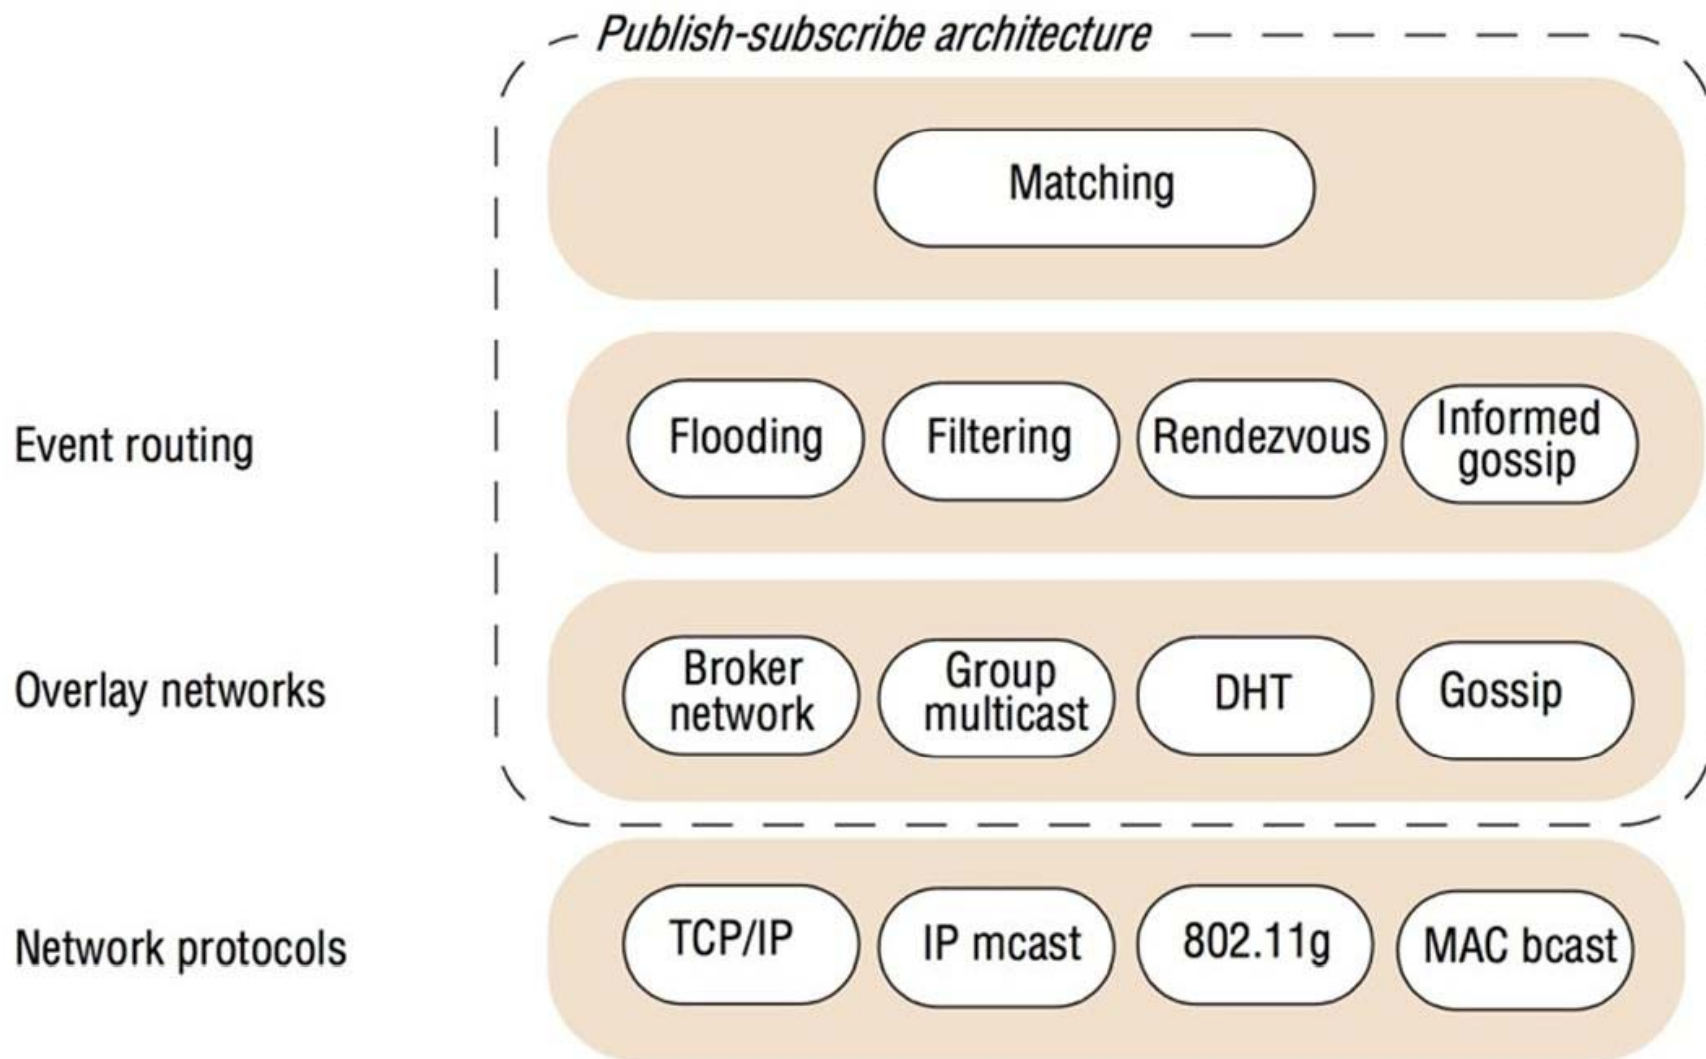
Publish-subscribe Paradigm

A

- ❑ (Event) Type-based subscription
 - Can organize events in hierarchy
- ❑ Channel-based subscription
 - Gets all events of that channel
- ❑ Content-based subscription
 - Filters events based on content



Architecture of Publish-Subscribe Systems

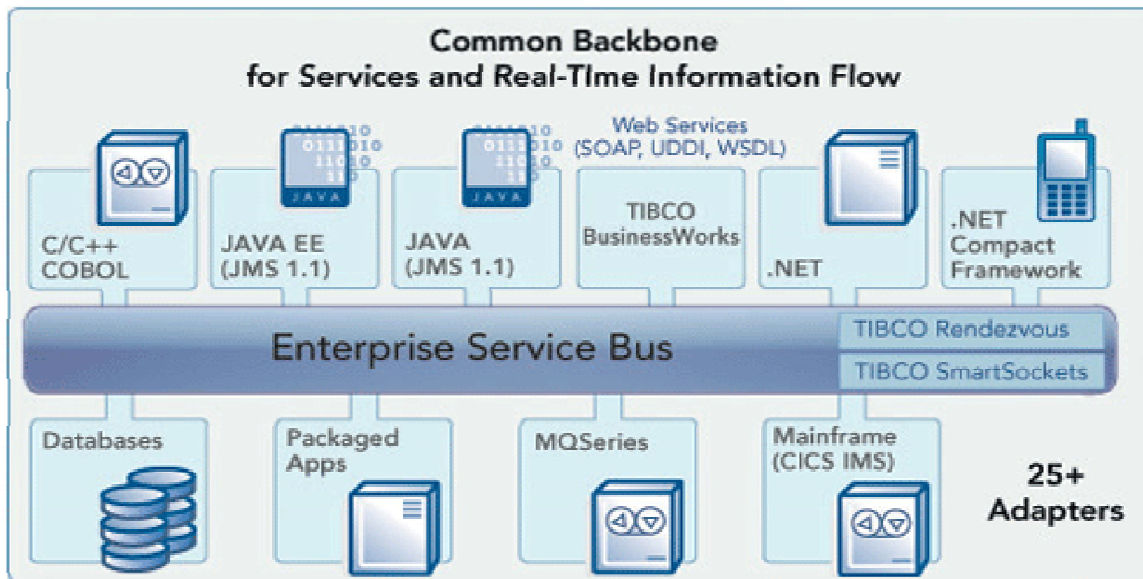


- ❑ Message Bus for Enterprise Application Integration
- ❑ Major design goals
 - Application-dependent communication system
 - Messages are self-describing
 - Processes should be referentially uncoupled
- ❑ Addresses are given as
 - Subject names
 - Inbox names
- ❑ Communication primitives
 - Send
 - SendRequest
 - SendReply

TIB: The Information Bus

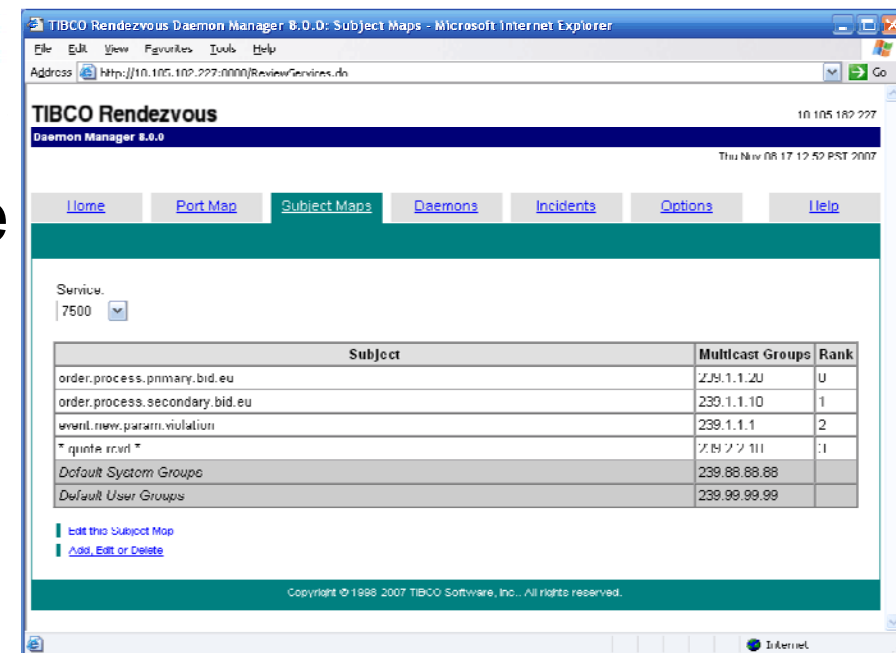
TIBCO – A TIB Instance

A



<https://www.tibco.com/products/tibco-rendezvous>

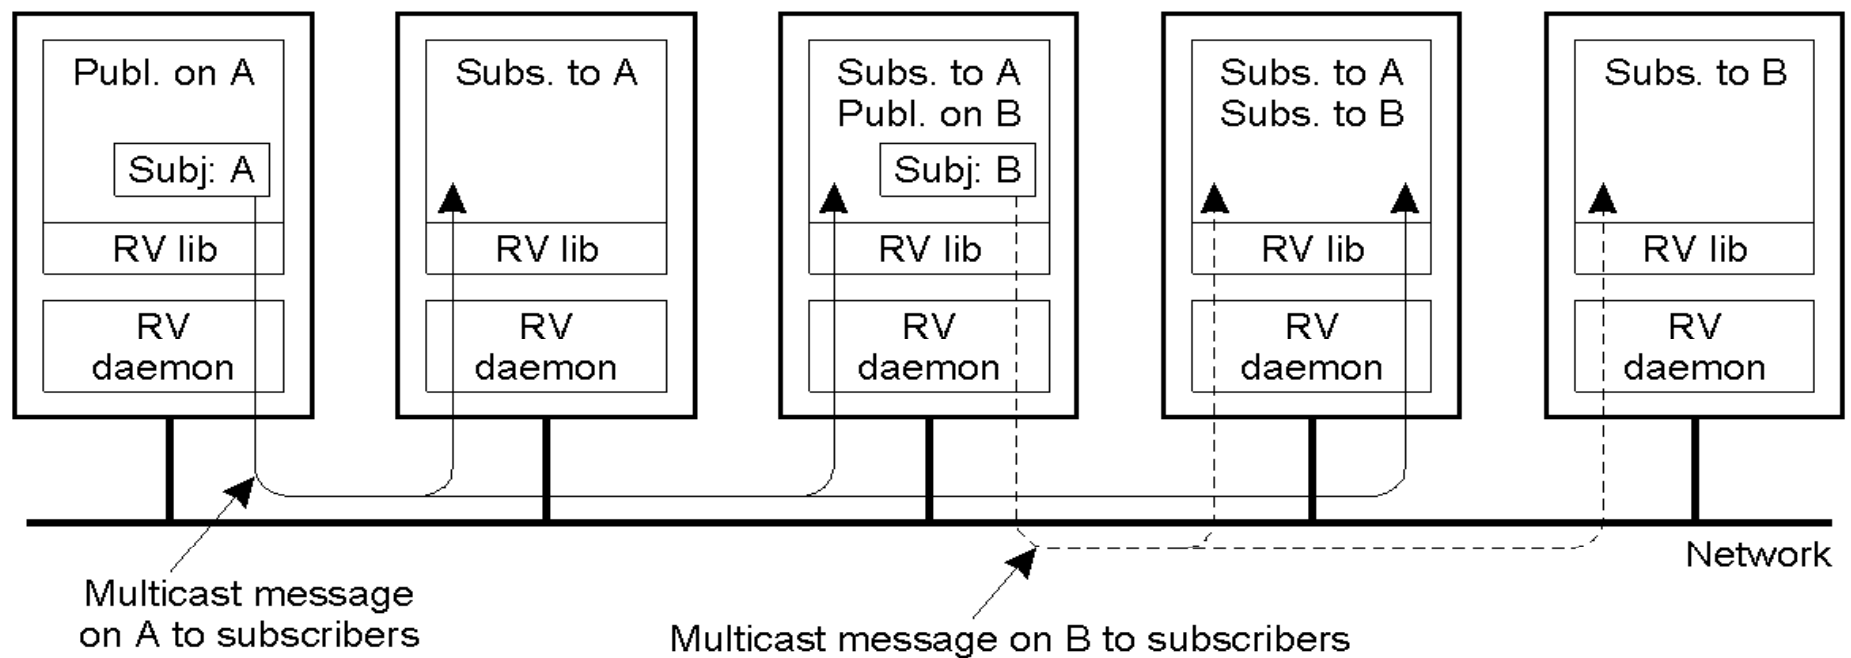
- ❑ Message-oriented Middleware
 - High-speed data distributions
 - Fully distributed daemon-based peer-to-peer architecture
 - No single point-of-failure



Coordination Model (1)

A

- Publish/subscribe system as in TIB/Rendezvous



Publ.: Publisher

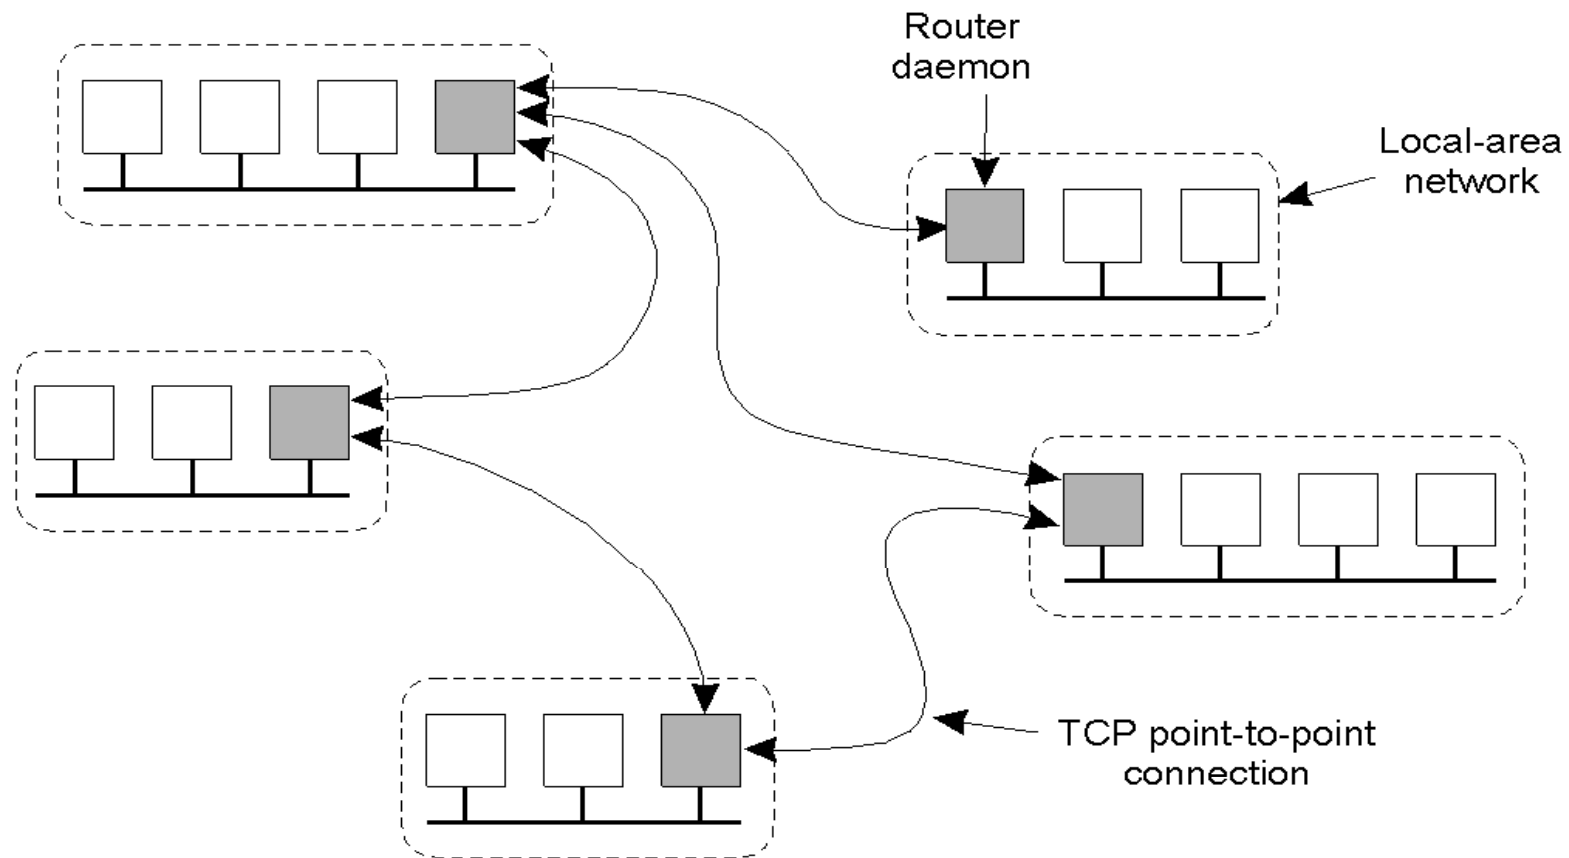
Subs.: Subscriber

RV: Rendezvous

Coordination Model (2)

A

- ❑ Overall architecture of a wide-area TIB/Rendezvous system



Basic Messaging

A

□ Example

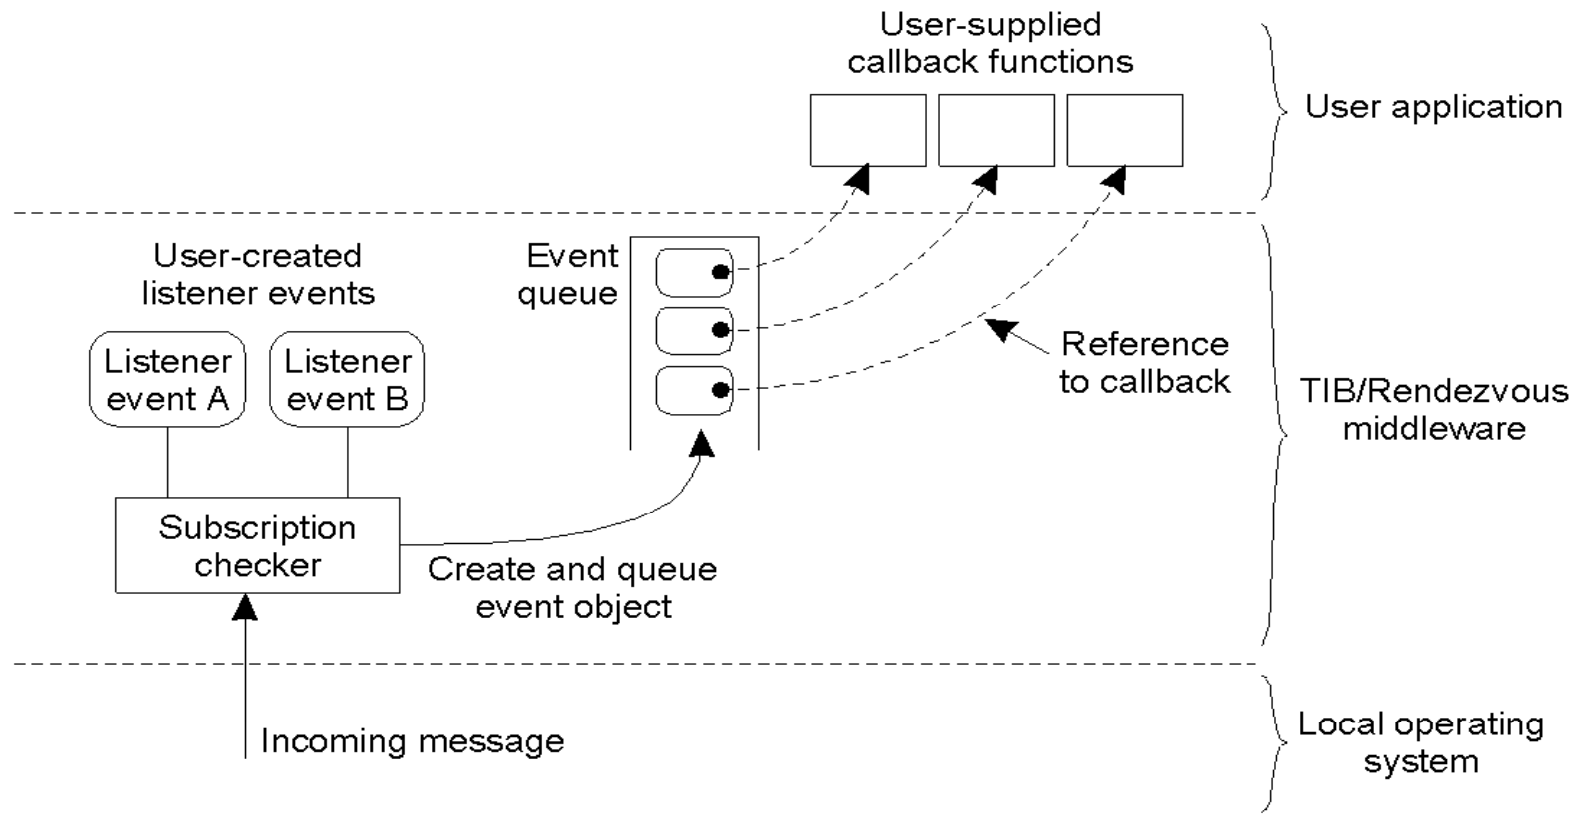
- Attributes of a TIB/Rendezvous message field

Attribute	Type	Description
Name	String	The name of the field, possibly NULL
ID	Integer	A message-unique field identifier
Size	Integer	The total size of the field (in bytes)
Count	Integer	The number of elements in the case of an array
Type	Constant	A constant indicating the type of data
Data	Any type	The actual data stored in a field

Events (1)

A

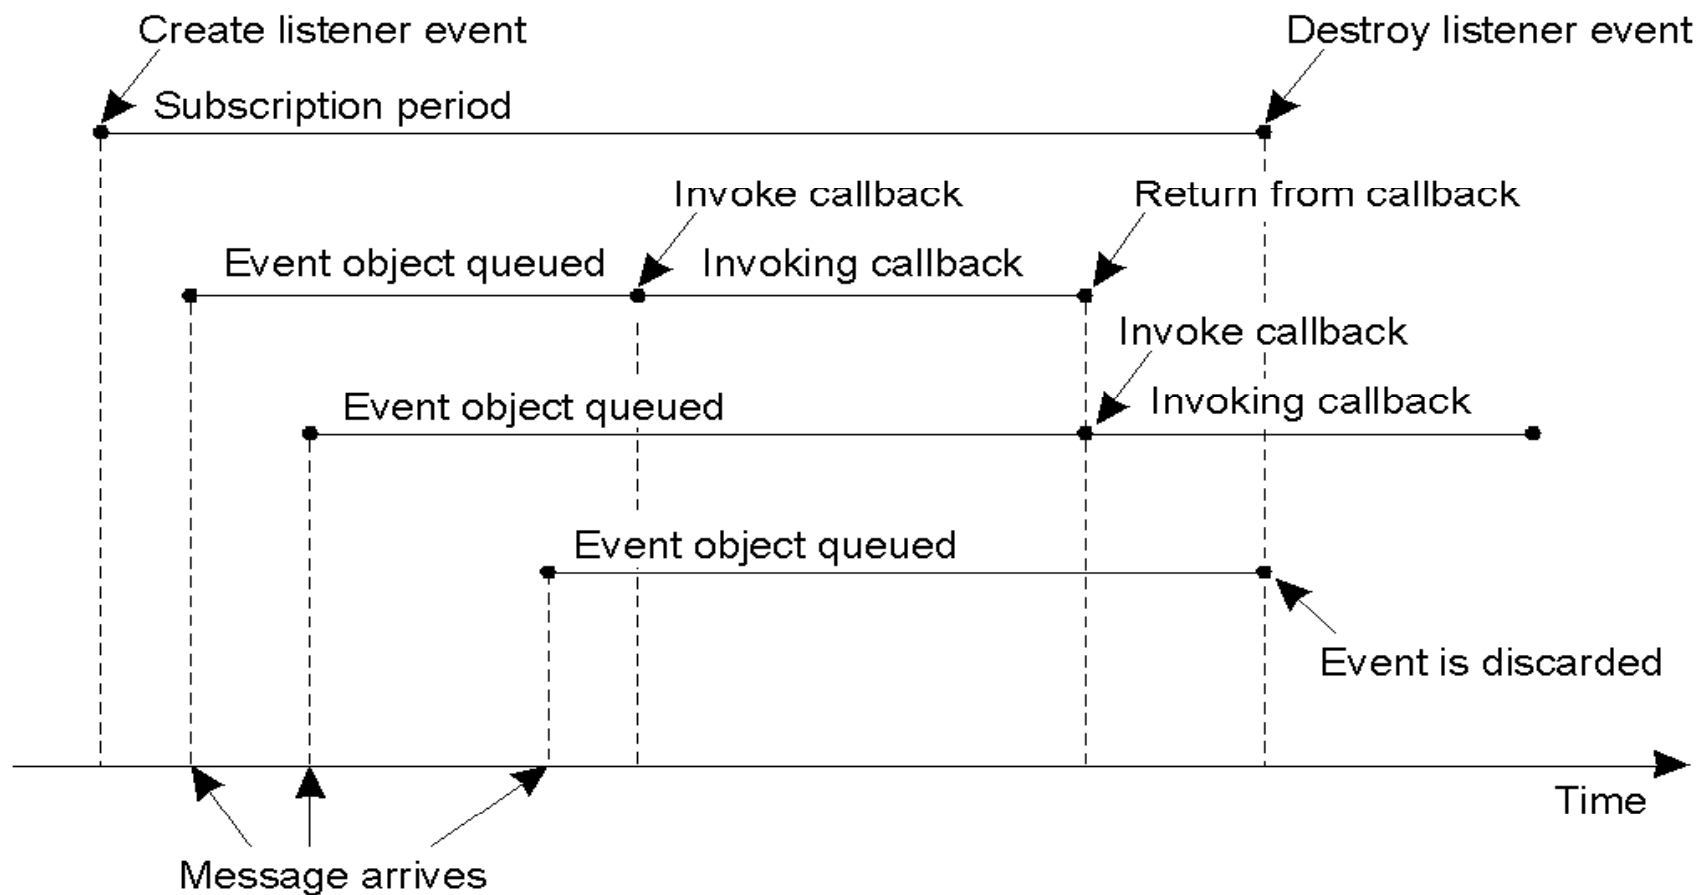
- Processing listener events for subscriptions in TIB/Rendezvous



Events (2)

A

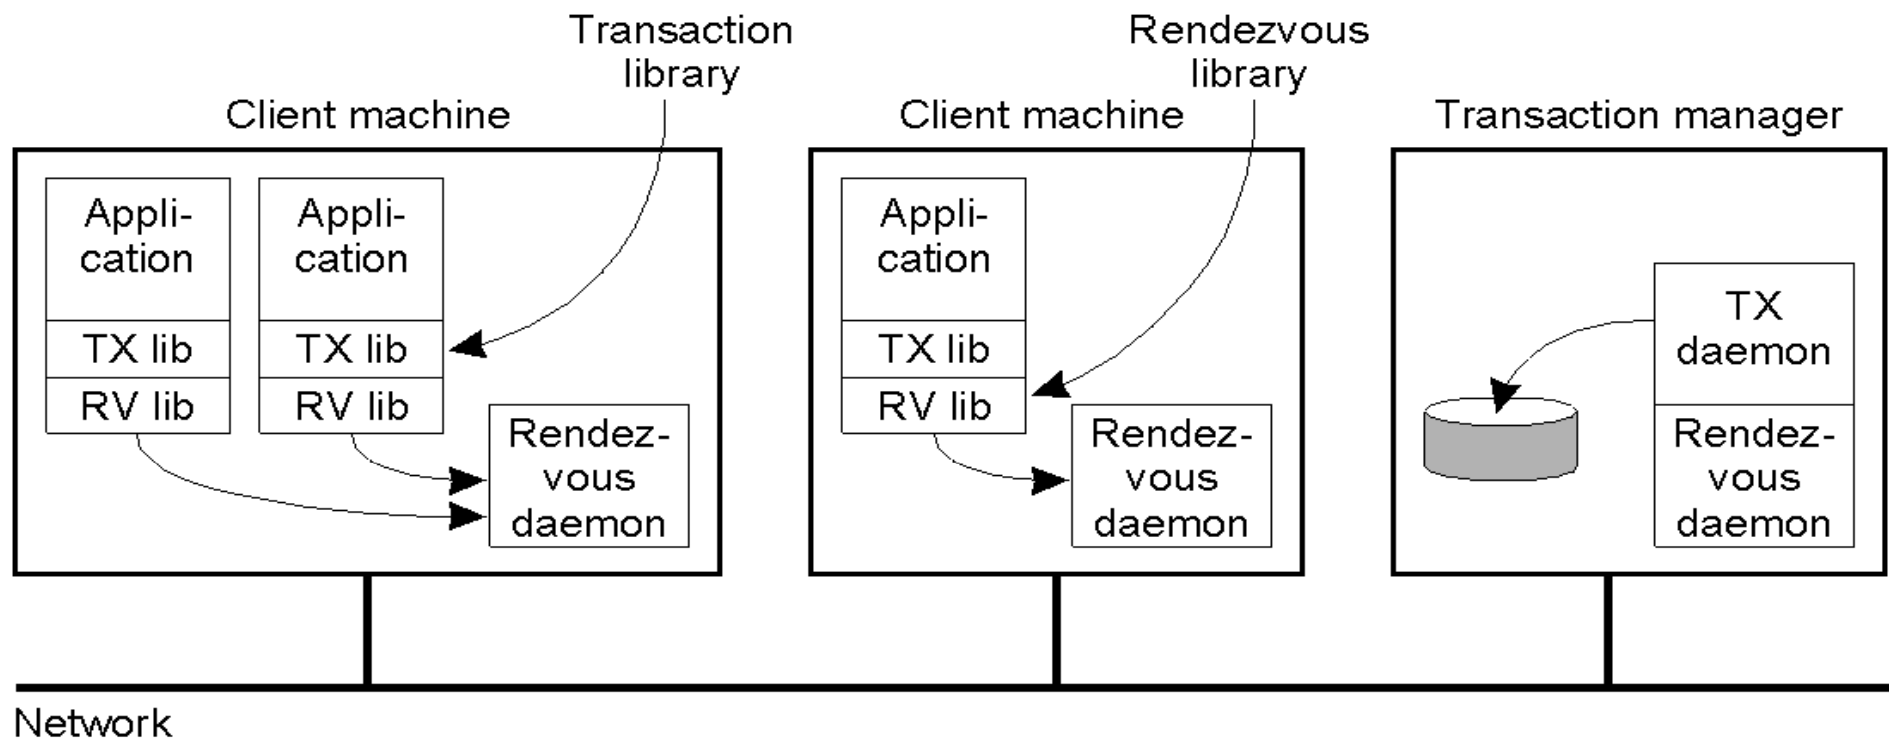
❑ Processing incoming messages in TIB/Rendezvous



Synchronization




A

- Organization of transactional messaging as a separate layer in TIB/Rendezvous




General Multicast

- a) A message is sent along a multicast tree
- b) A router will pass only a single NACK for each message
- c) A message is retransmitted only to receivers that have asked for it



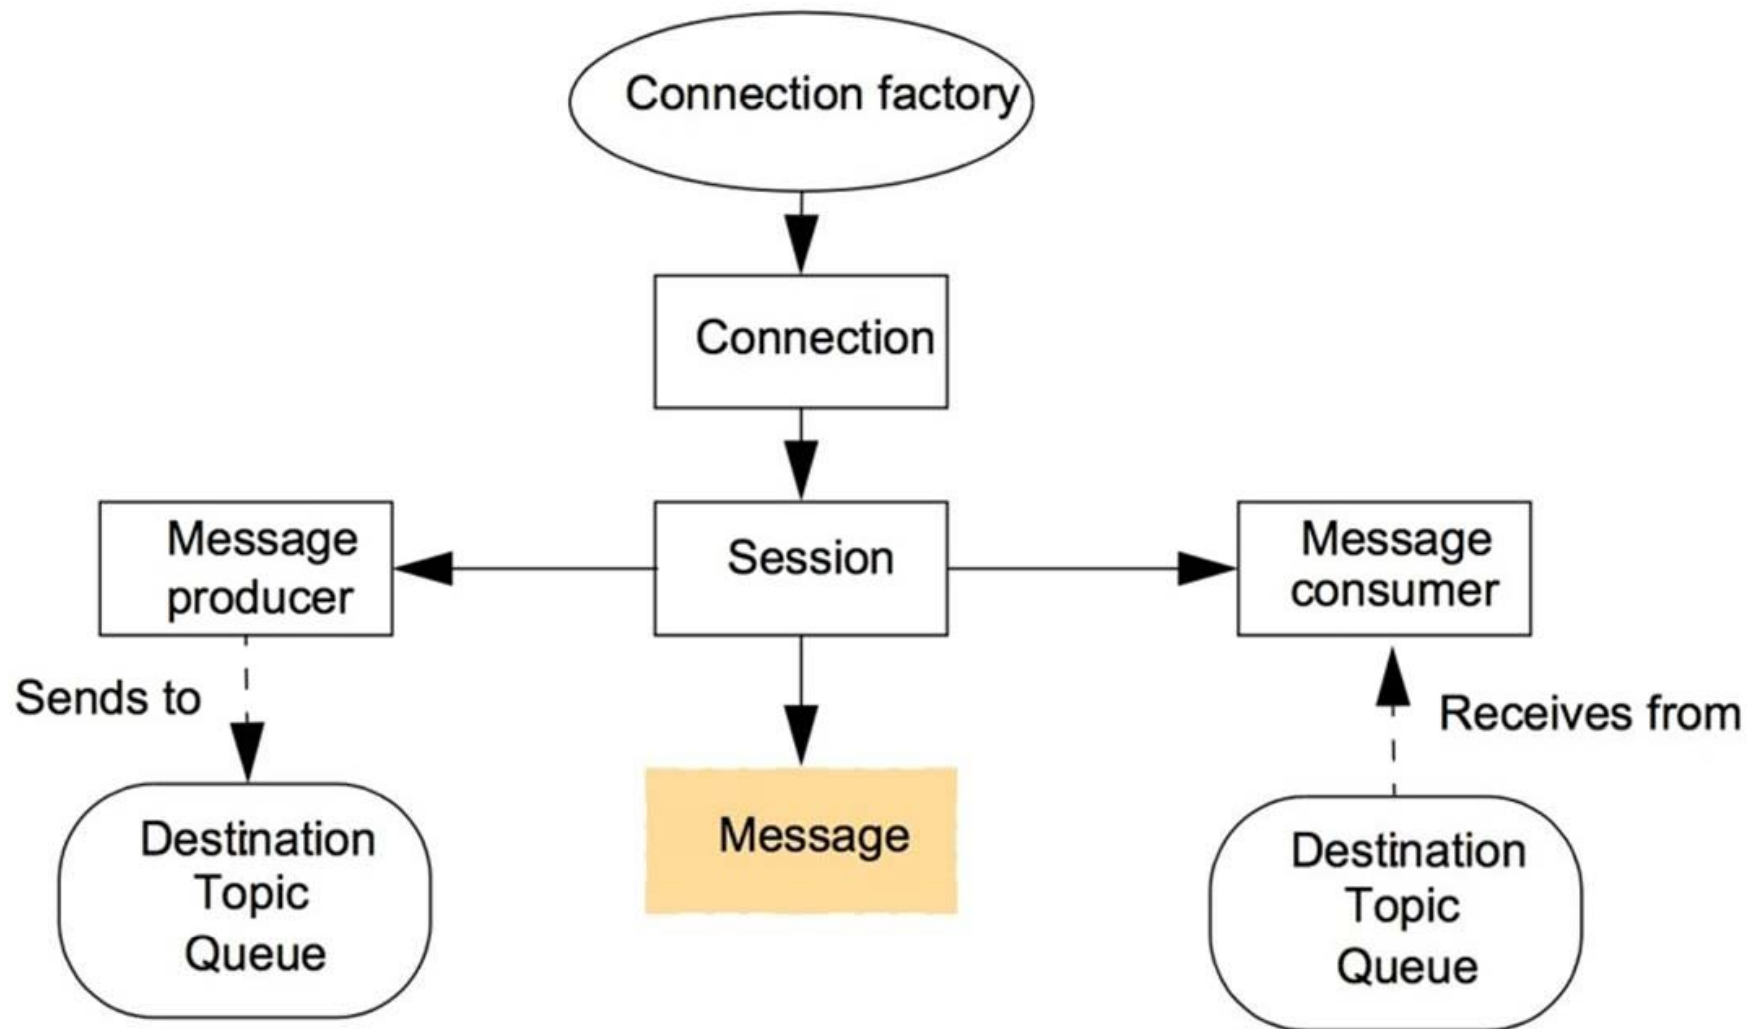
A Networked Topology: WebSphere MQ C

❑ MQ: Message Queue



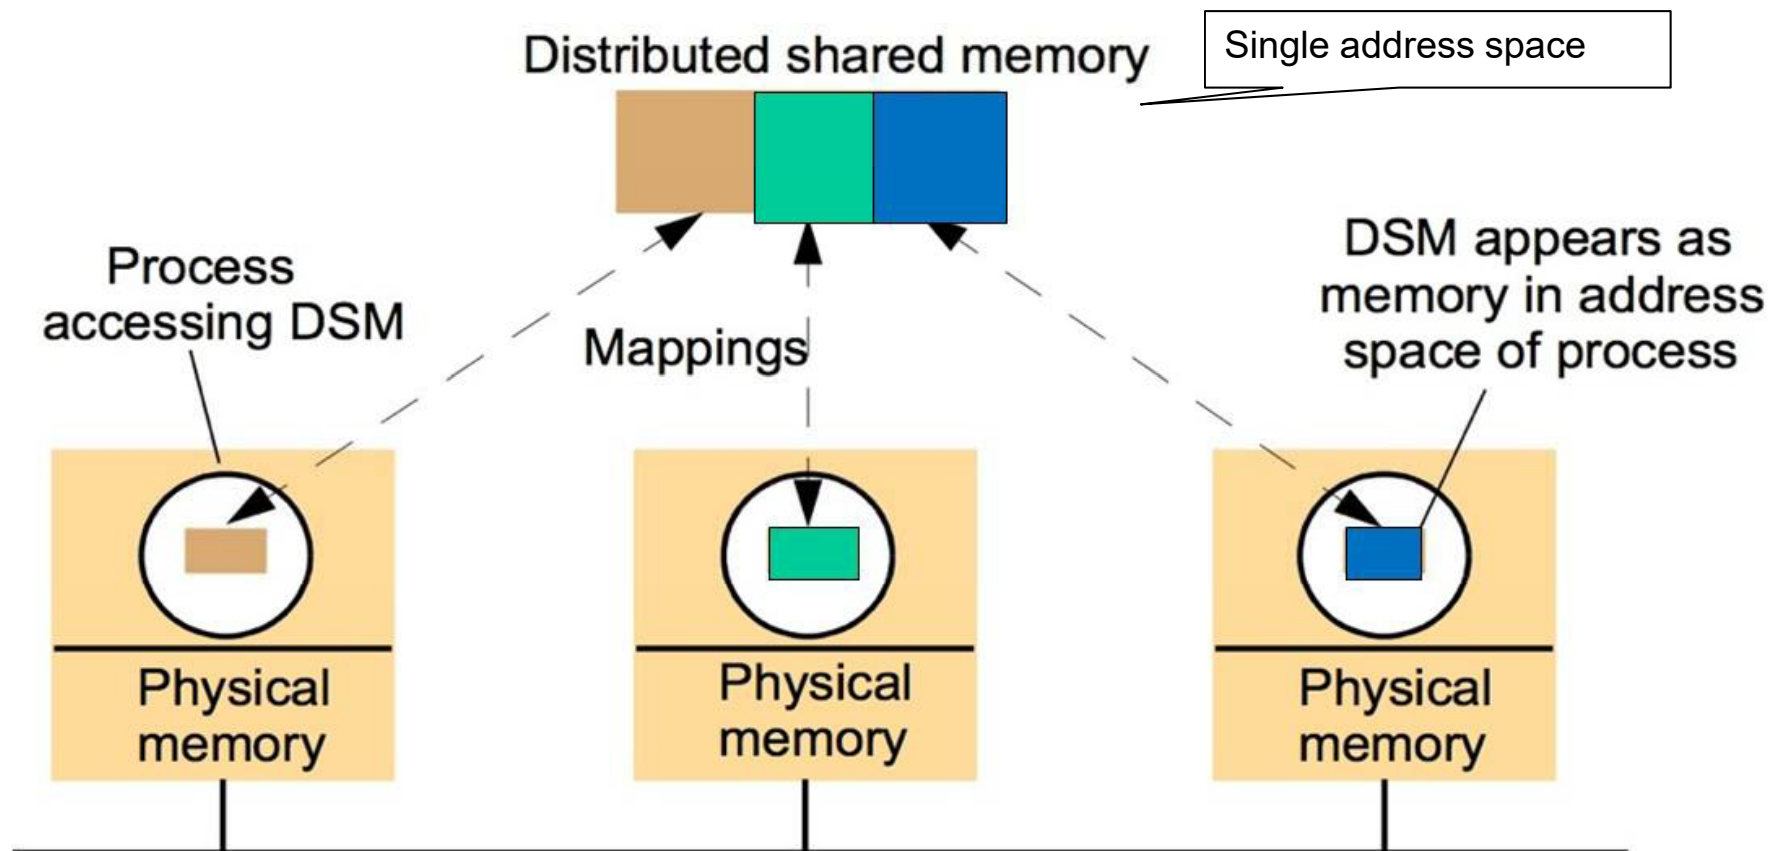
Java Messaging Service (JMS)

C



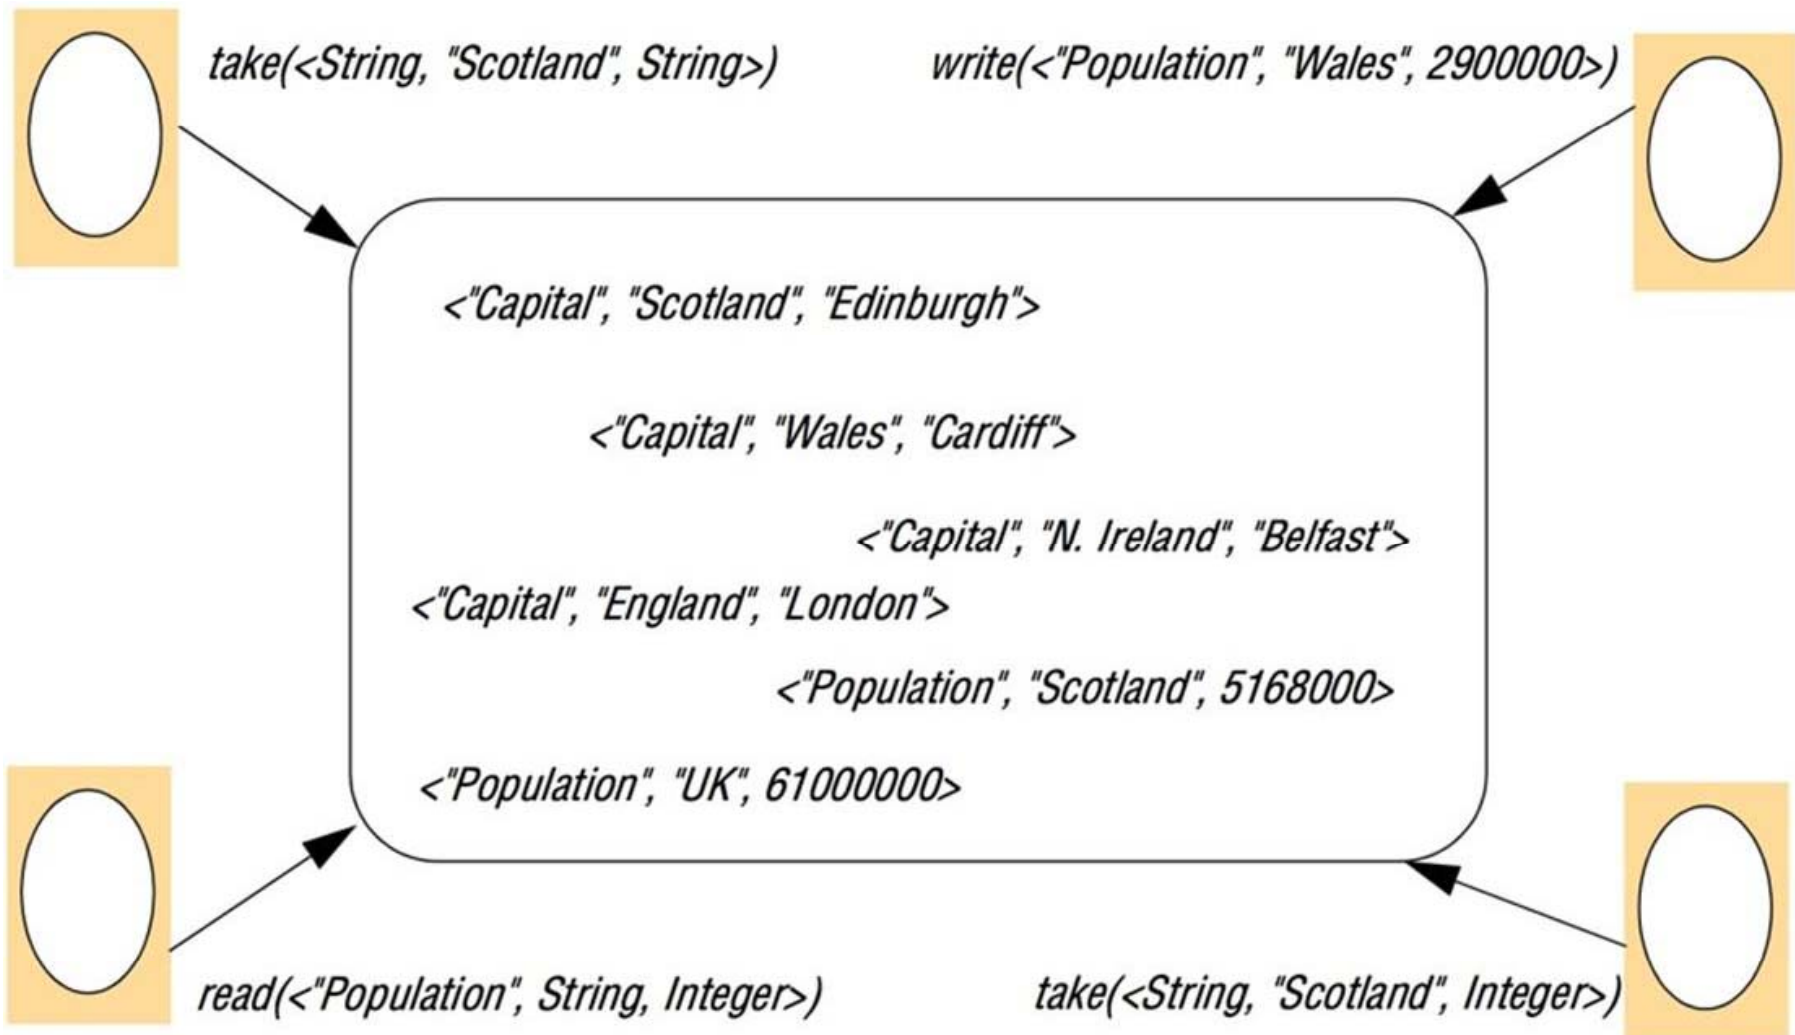
Distributed Shared Memory Abstraction C

- ❑ Indirect communication, based on bytes



Tuple Space Abstraction

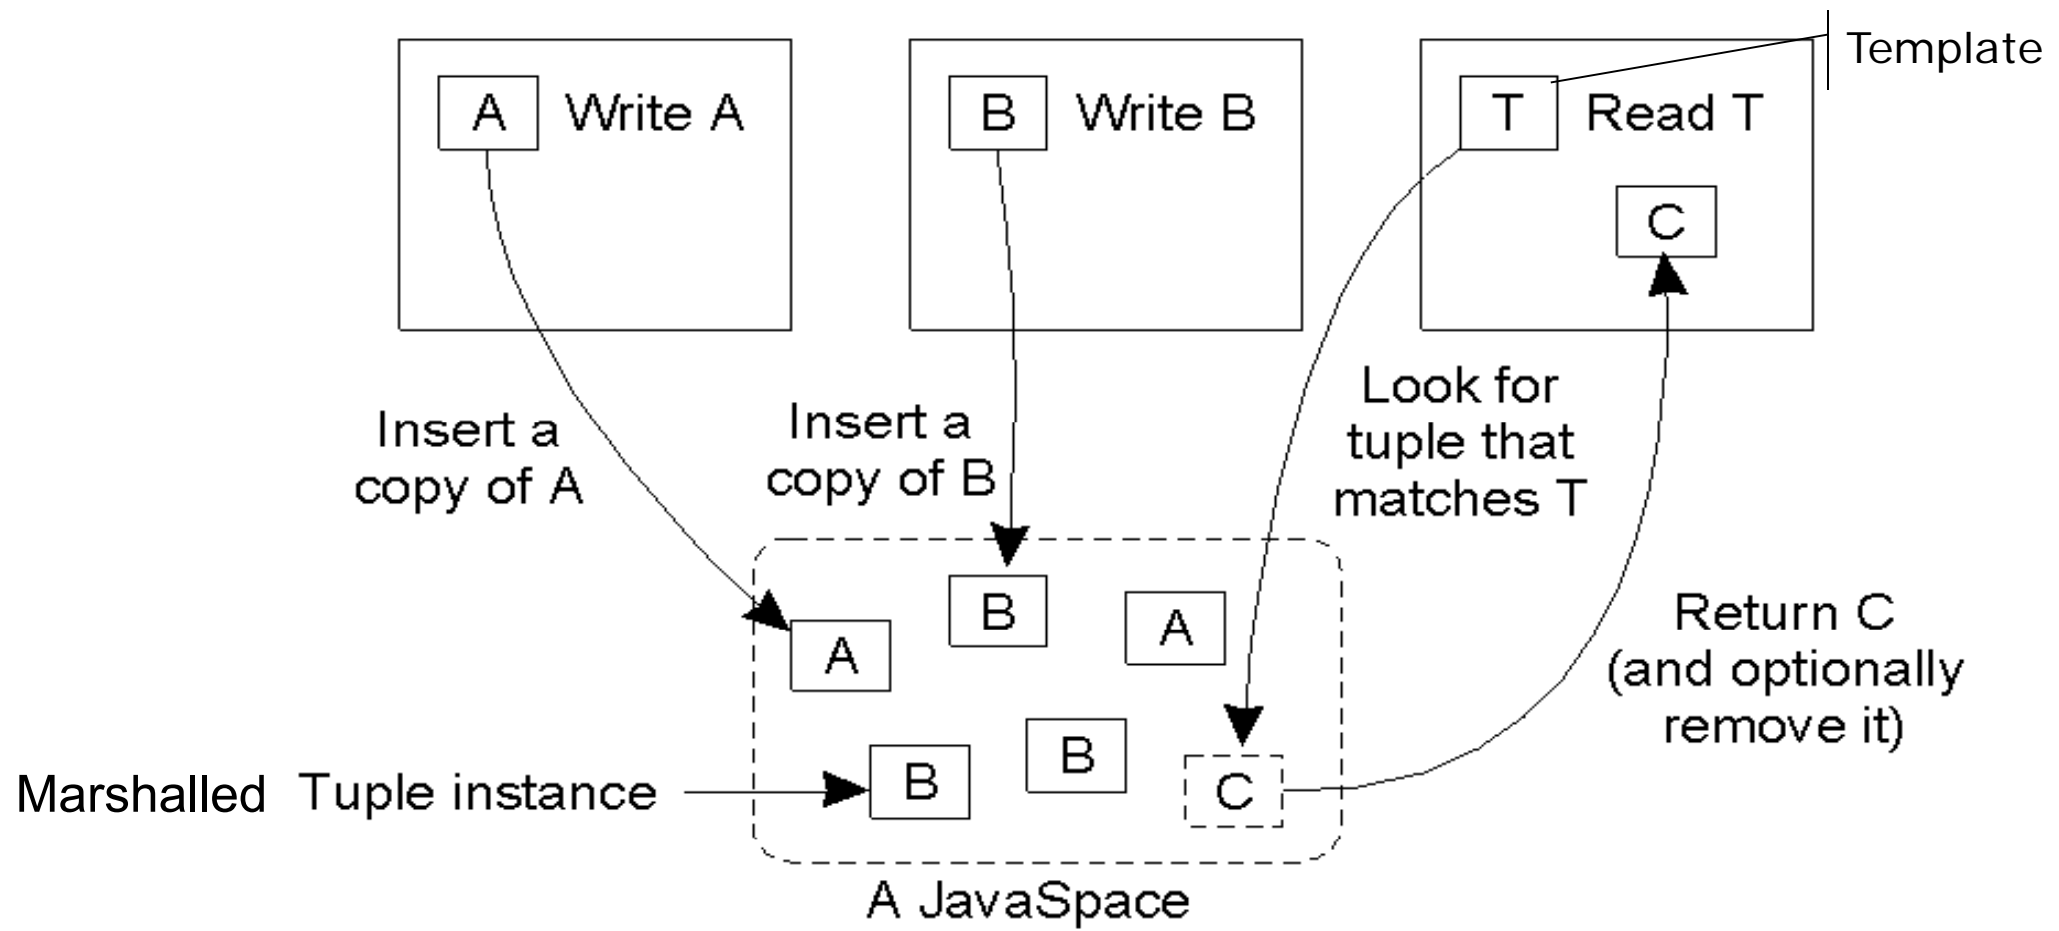
C



Jini Overview

C


□ General organization of a JavaSpace in Jini



Jini Architecture

C

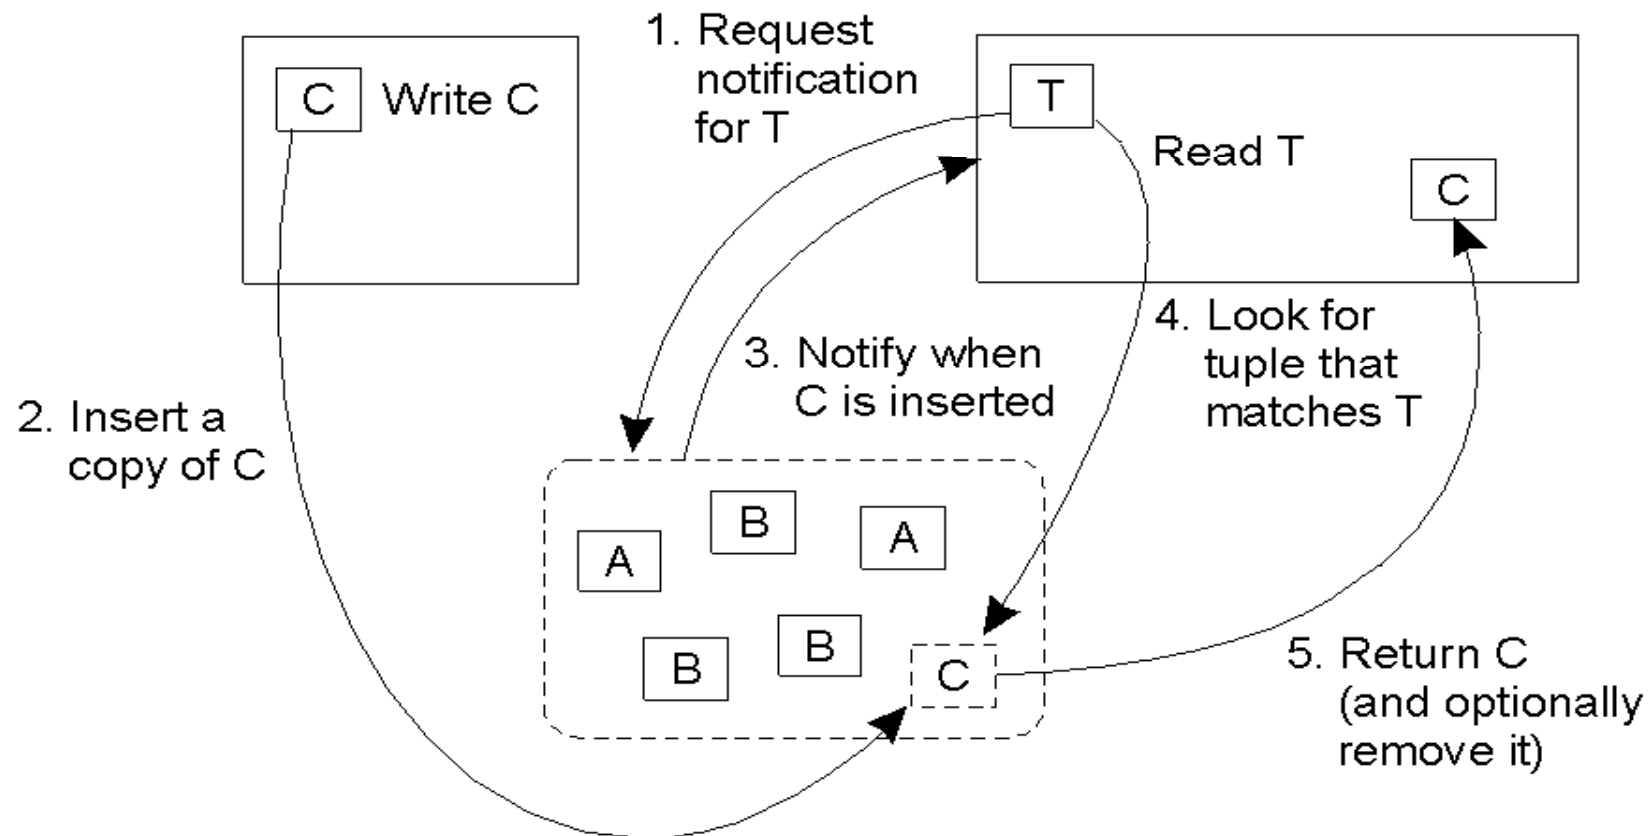
□ Layered architecture of Jini



Communication Events

C

- Using events in combination with a JavaSpace and Leases



Performance and Reliability Goals

C

- ❑ Servers with replication
- ❑ Caching on client side
- ❑ Zookeeper
 - Distributed coordination service
 - Publish-and-subscribe mechanism at hand (“watch”)
 - Hierarchical, scalable, atomicity, reliability
 - Sequential consistency
 - Uses Paxos to solve consensus

Example: Zookeeper in Hadoop Framework

